

# Parallel Computing with GPUs

## OpenMP

### Part 2 – Loops & Critical Sections



The  
University  
Of  
Sheffield.

Dr Paul Richmond

<http://paulrichmond.shef.ac.uk/teaching/COM4521/>



# This Lecture (learning objectives)

## ❑ Parallelising Loops

- ❑ Assign parallel section of code from loops to threads within OpenMP

## ❑ Critical Sections

- ❑ Identify the potential for race conditions in parallel code
- ❑ Examine a range of solutions for different race conditions



# OpenMP Syntax



## ❑ Parallel region directive

- ❑ `#pragma omp parallel [clause list] {structured block}`
- ❑ Spawns a number of parallel threads

## ❑ Clauses

- ❑ Are used to specify modifications to the parallel directive e.g.
  - ❑ Control scoping of variables in multiple threads
  - ❑ Dictate the number of parallel threads (example below)
  - ❑ Conditional parallelism

```
#pragma omp parallel num_threads(16)
{
    int thread = omp_get_thread_num();
    int max_threads = omp_get_max_threads();
    printf("Hello World (Thread %d of %d)\n", thread, max_threads);
}
```



# num\_threads()

- ❑ Without this clause `OMP_NUM_THREADS` will be used
  - ❑ This is an environment variable
  - ❑ Set to the number of cores (or hyperthreads) on your machine
  - ❑ This can be set globally by `omp_set_num_threads(int)`
  - ❑ Value can be queried by `int omp_get_num_threads();`
- ❑ `num_threads` takes precedence over the environment variable
- ❑ `num_threads()` does not guarantee that the number requested will be created
  - ❑ System limitations may prevent this
  - ❑ However: It almost always will

Application

Compiler

Environment

OpenMP Runtime

Platform threading model  
(e.g. Windows threading or pthreads)



# parallel for

- ❑ #pragma omp for
  - ❑ Assigns work units to the team
  - ❑ Divides loop iterations between threads
- ❑ For can be combined e.g. #pragma omp parallel for
  - ❑ Threads are spawned and then assigned to loop iterations



```
int n;
#pragma omp parallel for
for (n = 0; n < 8; n++){
    int thread = omp_get_thread_num();
    printf("Parallel thread %d \n", thread);
}
```

```
#pragma omp parallel
{
    int n;
    for (n = 0; n < 8; n++){
        int thread = omp_get_thread_num();
        printf("Parallel thread %d \n", thread);
    }
}
```

```
#pragma omp parallel
{
    int n;
    #pragma omp for
    for (n = 0; n < 8; n++){
        int thread = omp_get_thread_num();
        printf("Parallel thread %d \n", thread);
    }
}
```

Which is the odd one out?

# parallel for

- ❑ #pragma omp for
  - ❑ Assigns work units to the team
  - ❑ Divides loop iterations between threads
- ❑ For can be combined e.g. #pragma omp for
- ❑ Threads are spawned and then assigned

```
#pragma omp parallel
{
    int n;
    for (n = 0; n < 8; n++){
        int thread = omp_get_thread_num();
        printf("Parallel thread %d \n", thread);
    }
}
```

```
Parallel thread 0
Parallel thread 0
Parallel thread 0
Parallel thread 0
Parallel thread 0
Parallel thread 0
Parallel thread 0
Parallel thread 0
Parallel thread 2
Parallel thread 2
Parallel thread 2
Parallel thread 2
Parallel thread 2
Parallel thread 2
Parallel thread 2
Parallel thread 2
Parallel thread 2
Parallel thread 2
Parallel thread 5
Parallel thread 5
Parallel thread 5
Parallel thread 5
Parallel thread 5
Parallel thread 4
Parallel thread 4
Parallel thread 3
Parallel thread 3
Parallel thread 1
...
```



$$\cos(x) = \sum_{n=0}^{\infty} (-1)^n \frac{x^{2n}}{(2n)!}$$

$$\cos(x) = 1 - \frac{x^2}{2!} + \frac{x^4}{4!} - \frac{x^6}{6!} + \dots$$

# What is wrong with this code?

Consider a problem such as Taylor series expansion for cos function

$$\cos(x) = \sum_{n=0}^{\infty} (-1)^n \frac{x^{2n}}{(2n)!}$$

$$\cos(x) = 1 - \frac{x^2}{2!} + \frac{x^4}{4!} - \frac{x^6}{6!} + \dots$$

```
int n;
double result = 0.0;
double x = 1.0;

#pragma omp parallel for
for (n = 0; n < EXPANSION_STEPS; n++) {
    double r = pow(-1, n - 1) * pow(x, 2 * n - 1) / fac(2 * n);
    result -= r;
}

printf("Approximation of x is %f, value is %f\n", result, cos(x));
```





# Critical sections

❑ Consider a problem such as Taylor series expansion for *cos* function

$$\square \cos(x) = \sum_{n=0}^{\infty} (-1)^{n-1} \frac{x^{2n-1}}{(2n)!}$$

$$\square \cos(x) = 1 - \frac{x^2}{2!} + \frac{x^4}{4!} - \frac{x^6}{6!} \dots$$

```
int n;  
double result = 0.0;  
double x = 1.0;  
  
#pragma omp parallel for  
for (n = 0; n < EXPANSION_STEPS; n++) {  
    double r = pow(-1, n - 1) * pow(x, 2 * n - 1) / fac(2 * n);  
    result -= r;  
}  
  
printf("Approximation of x is %f, value is %f\n", result, cos(x));
```



**Race Condition:** Multiple threads try to write to the same value!  
(undefined behaviour and unpredictable results)



# Critical sections

❑ Consider a problem such as Taylor series expansion for *cos* function

$$\square \cos(x) = \sum_{n=0}^{\infty} (-1)^{n-1} \frac{x^{2n-1}}{(2n)!}$$

$$\square \cos(x) = 1 - \frac{x^2}{2!} + \frac{x^4}{4!} - \frac{x^6}{6!} \dots$$


```
int n;  
double result = 0.0;  
double x = 1.0;  
  
#pragma omp parallel for  
for (n = 0; n < EXPANSION_STEPS; n++) {  
    double r = pow(-1, n - 1) * pow(x, 2 * n - 1) / fac(2 * n);  
    #pragma omp critical  
    {  
        result -= r;  
    }  
}  
  
printf("Approximation of x is %f, value is %f\n", result, cos(x));
```

Solution: Define as a critical section



# Critical sections

- ❑ `#pragma omp critical [name]`
  - ❑ Ensures mutual exclusions when accessing a shared value
  - ❑ Prevents race conditions
  - ❑ A thread will wait until no other thread is executing a critical region (with the same name) before beginning
  - ❑ Unnamed critical regions map to the same unspecified name



# Atomics

- ❑ Atomic operations can be used to safely increment a shared numeric value
  - ❑ For example summation
  - ❑ Atomics only apply to the immediate assignment
- ❑ Atomics are usually faster than critical sections (benchmark to confirm)
  - ❑ Critical sections can be applied to general blocks of code (atomics can not)
- ❑ Example
  - ❑ Compute histogram of random values for a given range
  - ❑ Random is an `int` array of size `NUM_VALUES` with random value within `0 : RANGE`
  - ❑ Histogram is an `int` array of size `RANGE` with 0 values;

```
#pragma omp parallel
{
    int i;
    #pragma omp for
    for (i = 0; i < NUM_VALUES; i++){
        int value = randoms[i];
        #pragma omp atomic
        histogram[value]++;
    }
}
```



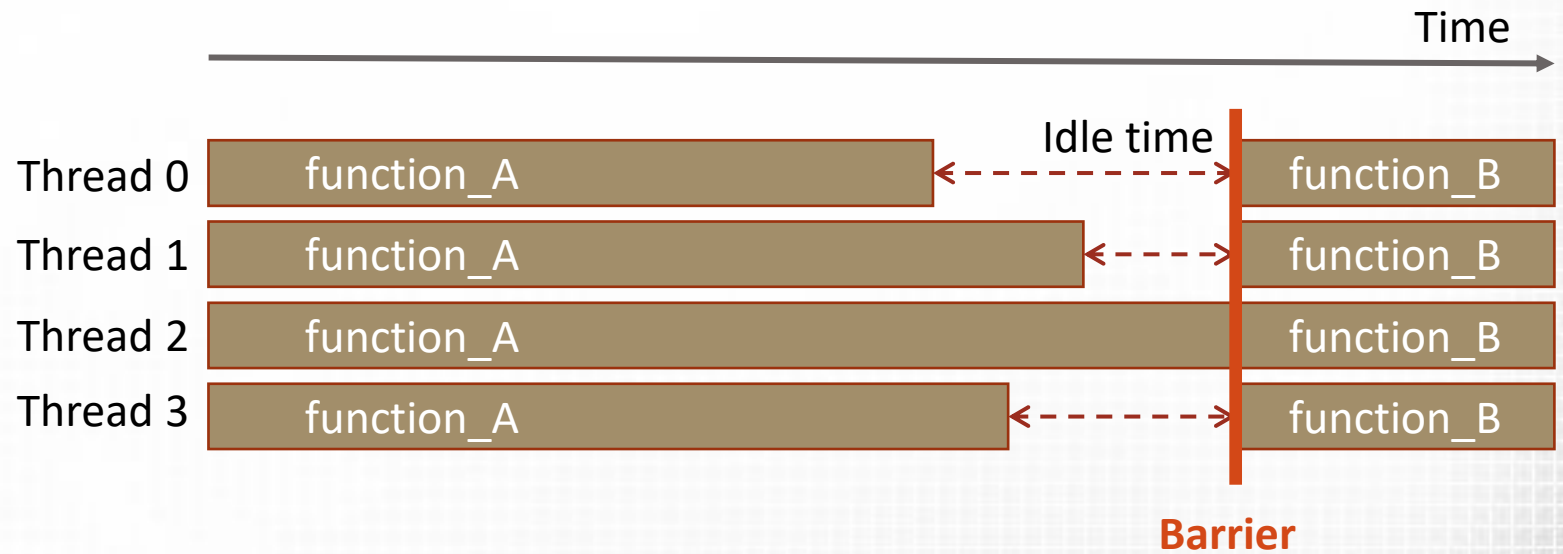




# Barriers

- ❑ `#pragma omp barrier`
  - ❑ Synchronises threads at a barrier point
  - ❑ Parallel regions have an implicit barrier
  - ❑ Can be used to ensure execution of particular code is complete
    - ❑ E.g. data read by `function_B`

```
#pragma omp parallel
{
    function_A()
    #pragma omp barrier
    function_B();
}
```



# Single and Master Sections



- ❑ `#pragma omp single { ... }`

- ❑ Used to ensure that only a single thread executes a region of a structured block

- ❑ Useful for I/O and initialisation

- ❑ First available thread will execute the defined region

- ❑ No control over which this is

- ❑ Will cause an implicit barrier (after structured block) unless a `nowait` clause is used

- ❑ E.g. `#pragma omp single nowait`

- ❑ `nowait` will remove an implied barrier and can also be applied to parallel for loops

- ❑ `#pragma omp master { ... }`

- ❑ Similar to `single` but will always use the primary/master thread

- ❑ Preferable to `single` (usually faster)

- ❑ Does not have an implicit barrier

# Master example

```
int t, r;
int local_histogram[THREADS][RANGE];

zero_histogram(local_histogram);

#pragma omp parallel num_threads(THREADS)
{
    int i;
    #pragma omp for
    for (i = 0; i < NUM_VALUES; i++){
        int value = randoms[i];
        local_histogram[omp_get_thread_num()][value]++;
    }
    #pragma omp barrier
    #pragma omp master
    for (t = 0; t < THREADS; t++){
        for (r = 0; r < RANGE; r++){
            histogram[r] += local_histogram[t][r];
        }
    }
}
```

Same result as the atomic version

Benchmark to understand performance!





# Summary

## ❑ Parallelising Loops

- ❑ Assign parallel section of code from loops to threads within OpenMP

## ❑ Critical Sections

- ❑ Identify the potential for race conditions in parallel code
- ❑ Examine a range of solutions for different race conditions

## ❑ Next Lecture: Scoping and Tasks

