

Parallel Computing with GPUs

Advanced OpenMP

Part 1 – Parallel Reduction



The
University
Of
Sheffield.

Dr Paul Richmond

<http://paulrichmond.shef.ac.uk/teaching/COM4521/>



This Lecture (learning objectives)

☐ Reduction

- ☐ Perform a parallel reduction using the reduction clause
- ☐ Recognise the limitations of the reduction functionality



❑ What do we need to look out for when considering applying OpenMP to this example?

```
void main() {  
    int i;  
    float vector[N];  
    float sum;  
  
    init_vector_values(vector);  
    sum = 0;  
  
    for (i = 0; i < N; i++) {  
        float v = some_func(vector[i]);  
        sum += v;  
    }  
    printf("Sum of values is %f\n", sum);  
}
```



Parallel Reduction

- ❑ A Reduction is the combination of local copies of a variable into a single copy
 - ❑ Consider a case where we want to sum the values of a function operating on a vector of values;

```
void main() {  
    int i;  
    float vector[N];  
    float sum;  
  
    init_vector_values(vector);  
    sum = 0;  
  
    for (i = 0; i < N; i++) {  
        float v = some_func(vector[i]);  
        sum += v;  
    }  
    printf("Sum of values is %f\n", sum);  
}
```

Candidate for
parallel
reduction...



Reduction clause

```
void main() {  
    int i;  
    float vector[N];  
    float sum;  
  
    init_vector_values(vector);  
    sum = 0;  
  
    #pragma omp parallel for reduction(+: sum);  
    for (i = 0; i < N; i++) {  
        float v = some_func(vector[i]);  
        sum += v;  
    }  
    printf("Sum of values is %f\n", sum);  
}
```

Without reduction we would need a critical section to update the shared variable!



OpenMP Reduction

- ❑ Reduction is supported with the reduction clause which requires a reduction variable
 - ❑ E.g. `#pragma omp parallel reduction(+: sum_variable) {...}`
 - ❑ Reduction variable is implicitly private to other threads
- ❑ OpenMP implements this **in parallel** by;
 - ❑ Creating a local (private) copy of the (shared) reduction variable
 - ❑ Combining (merging) local copies of the variable at the end of the structured block
 - ❑ Saving the reduced value to the shared variable in the master thread.
- ❑ Reduction operators are `+`, `-`, `*`, `&`, `|`, `&&` and `||`
 - ❑ `&`: bitwise and
 - ❑ `|`: bitwise or
 - ❑ `&&`: logical and
 - ❑ `||`: logical or



Summary

❑ Reduction

- ❑ Perform a parallel reduction using the reduction clause
- ❑ Recognise the limitations of the reduction functionality

❑ Next Lecture: Scheduling

