

Parallel Computing with GPUs

Advanced OpenMP Part 2 – Scheduling



The
University
Of
Sheffield.

Dr Paul Richmond

<http://paulrichmond.shef.ac.uk/teaching/COM4521/>



This Lecture (learning objectives)

☐ Scheduling

- ☐ Compare and contrast different scheduling approaches to understand the benefits and limitations of each
- ☐ Identify how scheduling parameters may impact cache utilisation





Scheduling clause

- ❑ OpenMP by default uses static scheduling
 - ❑ Static: schedule is determined at compile time
 - ❑ `schedule (static)`
- ❑ In general: `schedule (type [, chunk size])`
 - ❑ `type=static`: Iterations assigned to threads before execution (preferably at compile time)
 - ❑ *`type=dynamic`: iterations are assigned to threads as they become available*
 - ❑ *`type=guided`: iterations are assigned to threads as they become available (with reducing chunk size)*
 - ❑ `type=auto`: compiler and runtime determine the schedule
 - ❑ `type=runtime`: schedule is determined at runtime by env variable

What would be a use case where static scheduling is a bad choice?



Static scheduling chunk size

❑ chunk size

- ❑ Refers to the amount of work assigned to each thread
- ❑ By default chunk size is to divide the work by the number of threads
 - ❑ Low overhead (no going back for more work)
 - ❑ Not good for uneven workloads
 - ❑ E.g. consider our last lectures Taylor series example (updated to use reduction)

```
int n;  
double result = 0.0;  
double x = 1.0;  
  
#pragma omp parallel for reduction(-: result)  
for (n = 0; n < EXPANSION_STEPS; n++){  
    double r = pow(-1, n - 1) * pow(x, 2 * n - 1) / fac(2 * n);  
    result -= r;  
}  
  
printf("Approximation is %f, value is %f\n", result, cos(x));
```

Recursive uneven workload



Scheduling Workload




```
long long int fac(int n)
{
    if (n == 0)
        return 1;
    else
        return(n * fac(n - 1));
}
```

❑ Uneven workload amongst threads

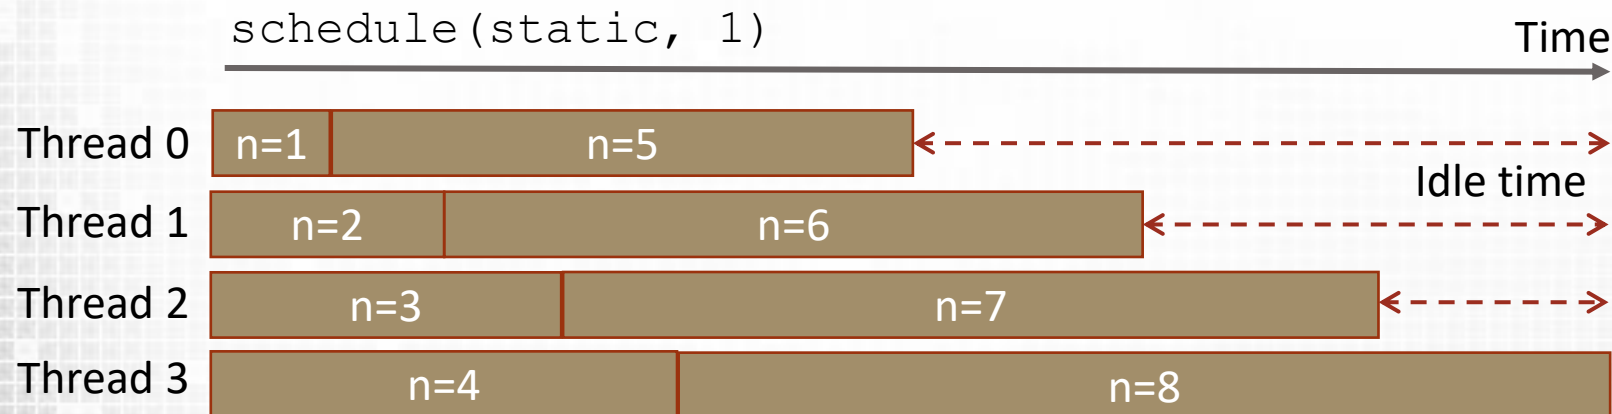
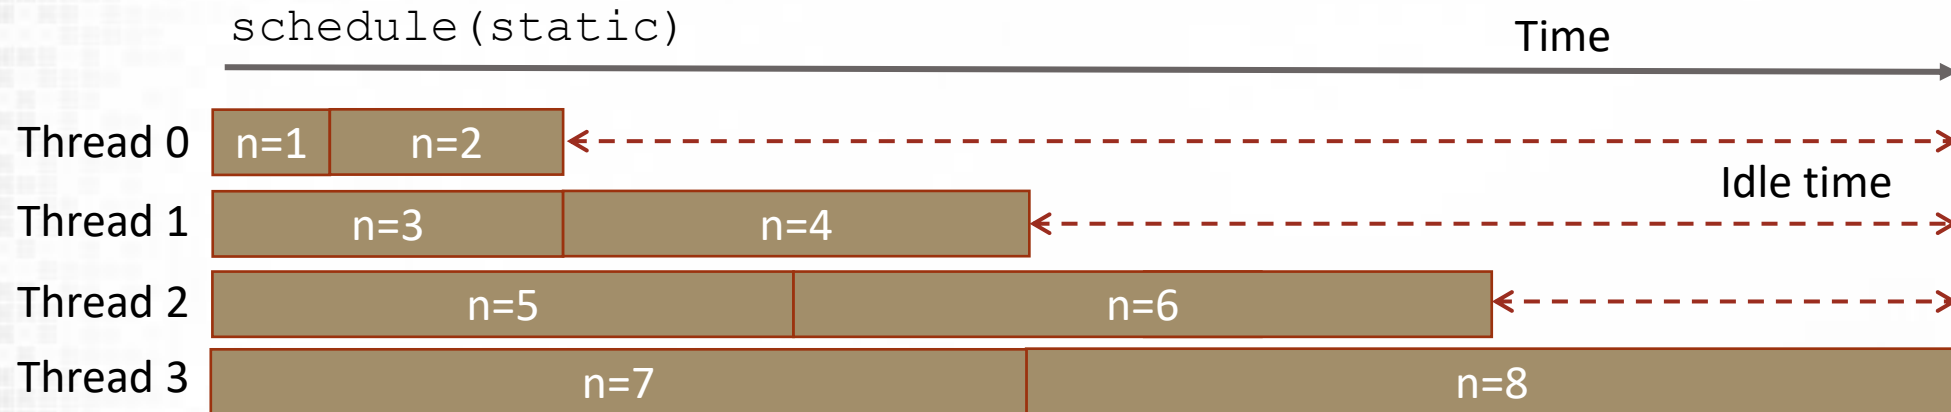
❑ Increase in n leads to increased computation

❑ E.g. `EXPANSION_STEPS=8, num_threads(4), schedule(static)`



Cyclic Scheduling

- ❑ It would be better to partition the workload more evenly
 - ❑ E.g. Cyclic scheduling via chunk size



Cyclic Scheduling

```
#pragma omp for num_threads(4)
for (i = 0; i < 16; i++)
```

`schedule(static, 1)`

Thread 0	0	4	8	12
Thread 1	1	5	9	13
Thread 2	2	6	10	14
Thread 3	3	7	11	15

`schedule(static, 2)`

Thread 0	0	1	8	9
Thread 1	2	3	10	11
Thread 2	4	5	12	13
Thread 3	6	7	14	15

`schedule(static, 4)`

Thread 0	0	1	2	3
Thread 1	4	5	6	7
Thread 2	8	9	10	11
Thread 3	12	13	14	15

Default case

- ❑ Default chunk size is $n/\text{threads}$
- ❑ where n is the number of iterations



Dynamic and Guided Scheduling

☐ Dynamic

- ☐ **Iterations are broken down by chunk size**
- ☐ Threads request chunks of work from a runtime queue when they are free
- ☐ Default chunk size is 1

☐ Guided

- ☐ **Chunks of the workload grow exponentially smaller**
- ☐ Threads request chunks of work from a runtime queue when they are free
- ☐ Chunk size is the size which the workloads decrease to
 - ☐ with the exception of last chunk which may have remainder

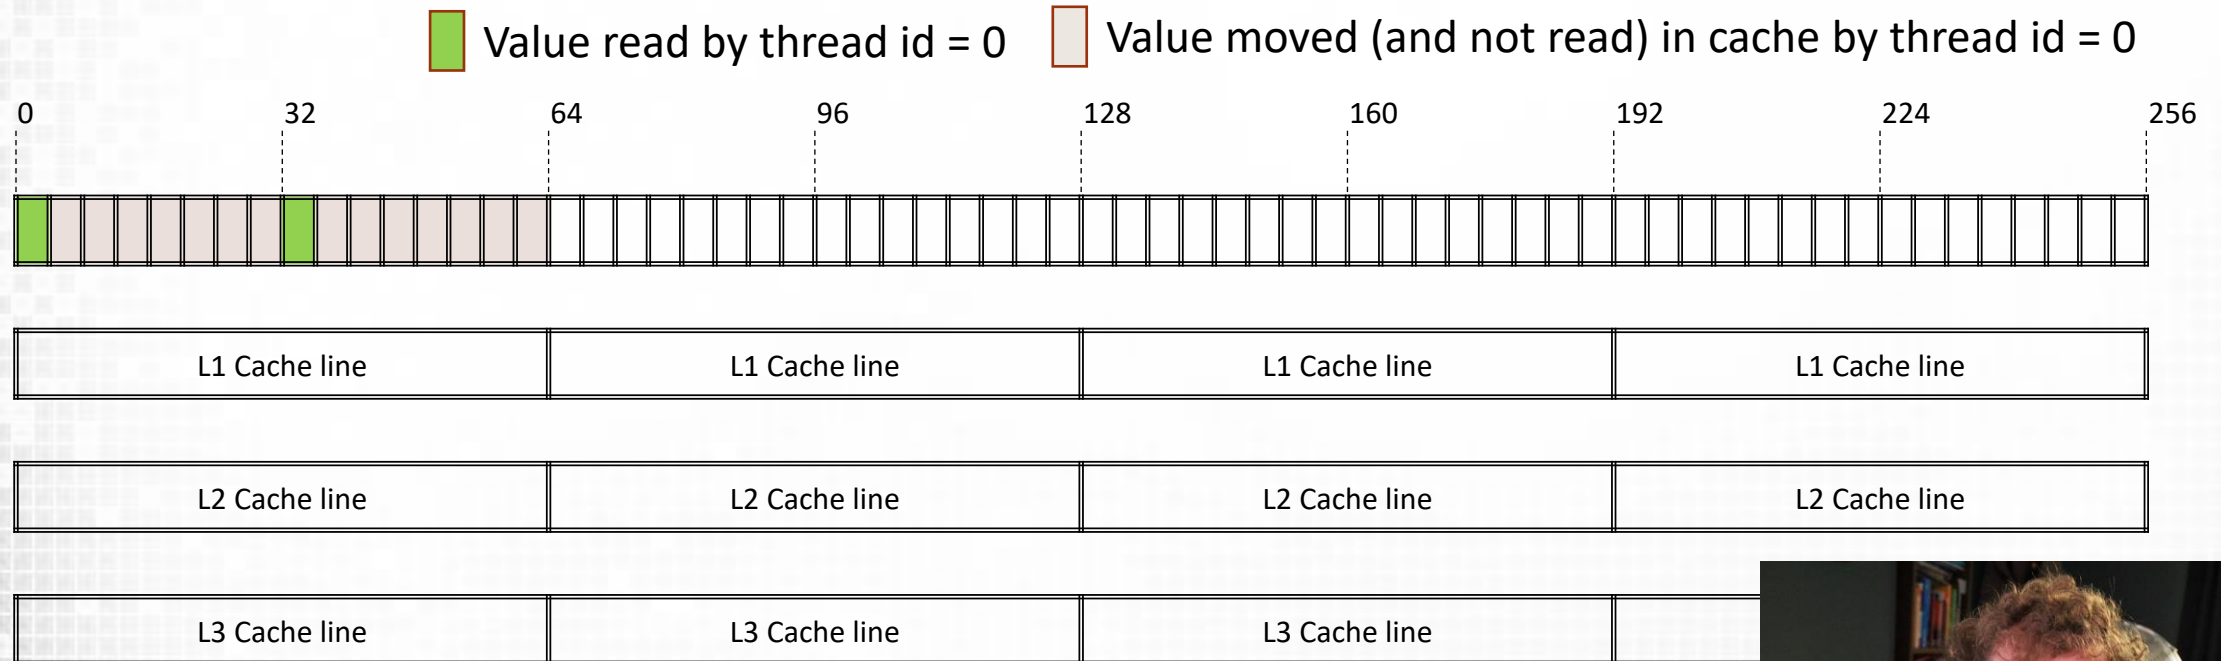
☐ Both

- ☐ Requesting work dynamically creates overhead
 - ☐ Not well suited if iterations are balanced
- ☐ Overhead vs. imbalance: How do I decide which is best?
 - ☐ **Benchmark all to find the best solution**



Cache Efficiency

```
#pragma omp parallel for schedule(static,1) num_threads(8)
  for (int i=0; i<64; i++) {
    something(array[i]);
  }
```




❑ Chunk size may effect cache utilisation



False Sharing


- ❑ Changing a single value causes a whole cache line to be invalid
 - ❑ Invalid lines must be re-cached
 - ❑ Chunk size case effect the amount of times a line is invalid

```
#pragma omp parallel for schedule(static,1)
for (int i=0; i<64; i++) {
    array[i]++;
}
```



False Sharing (worked example)

```
#pragma omp parallel for schedule(static,1) num_threads(8)
for (int i=0; i<64; i++) {
    array[i]++;
}
```



Summary

❑ Scheduling

- ❑ Compare and contrast different scheduling approaches to understand the benefits and limitations of each
- ❑ Identify how scheduling parameters may impact cache utilisation

❑ Next Lecture: Nesting Loops and OpenMP Summary

