

Parallel Computing with GPUs

Advanced OpenMP

Part 3 – Nesting and Summary



The
University
Of
Sheffield.

Dr Paul Richmond

<http://paulrichmond.shef.ac.uk/teaching/COM4521/>



This Lecture (learning objectives)

☐ Nesting

- ☐ Operate on nested loops using OpenMP
- ☐ Compare the performance implications of different approaches for nesting

☐ Summary

- ☐ Classify permitted use of the various OpenMP clauses



Nesting

- ❑ Consider the following example...
- ❑ How should we parallelise this example?

```
for (i = 0; i < OUTER_LOOPS; i++) {  
    for (j = 0; j < INNER_LOOPS; j++) {  
        printf("Hello World (Thread %d)\n", omp_get_thread_num());  
    }  
}
```



Nesting

- ❑ Consider the following example...
- ❑ How should we parallelise this example?

```
#pragma omp parallel for
for (i = 0; i < OUTER_LOOPS; i++){
    for (j = 0; j < INNER_LOOPS; j++){
        printf("Hello World (Thread %d)\n", omp_get_thread_num());
    }
}
```

- ❑ What if $\text{OUTER_LOOPS} \ll \text{number of threads}$
- ❑ E.g. $\text{OUTER_LOOPS} = 2$



Nesting

❑ We can use parallel nesting

❑ Nesting is turned off by default so we must use `omp_set_nested()`

❑ When inner loop is met each outer thread creates a new team of threads

❑ Allows us to expose higher levels of parallelism

❑ *Only useful when outer loop does not expose enough parallelism*


```
omp_set_nested(1);

#pragma omp parallel for
for (i = 0; i < OUTER_LOOPS; i++){
    int outer_thread = omp_get_thread_num();
    #pragma omp parallel for
    for (j = 0; j < INNER_LOOPS; j++){
        int inner_thread = omp_get_thread_num();
        printf("Hello World (i T=%d j T=%d)\n", outer_thread, inner_thread);
    }
}
```

```
Hello World (i T=0 j T=0)
Hello World (i T=0 j T=1)
Hello World (i T=0 j T=3)
Hello World (i T=1 j T=2)
Hello World (i T=1 j T=1)
Hello World (i T=1 j T=0)
Hello World (i T=0 j T=2)
Hello World (i T=1 j T=3)
```

Nesting Fork and Join

- ❑ Every parallel directive creates a fork (new team)
 - ❑ In this case `omp parallel` is used to fork a new parallel region



Collapse

- ❑ Only available in OpenMP 3.0 and later (not VS2017)
 - ❑ Can automatically collapse multiple loops
 - ❑ Loops must not have statements or expressions between them

```
#pragma omp parallel for collapse(2)
for (i = 0; i < OUTER_LOOPS; i++){
    for (j = 0; j < INNER_LOOPS; j++){
        int thread = omp_get_thread_num();
        printf("Hello World (T=%d)\n", thread);
    }
}
```

Work around...

```
#pragma omp parallel for
for (i = 0; i < OUTER_LOOPS* INNER_LOOPS; i++){
    int thread = omp_get_thread_num();
    printf("Hello World (T=%d)\n", thread);
}
```




Clauses usage summary

| Clause | Directive: #pragma omp ... | | | | | |
|--------------|----------------------------|-----|----------|--------|-----------------|----------------------|
| | parallel | for | sections | single | parallel for | parallel sections |
| if | | | | | | |
| private | | | | | | |
| shared | | | | | | |
| default | | | | | | |
| firstprivate | | | | | | |
| lastprivate | | | | | | |
| reduction | | | | | | |
| schedule | | | | | | |
| nowait | | | | | | |



Performance

- ❑ Remember ideas for general C performance
 - ❑ Have good data locality (good cache usage)
 - ❑ Combine loops where possible
- ❑ Additional performance criteria
 - ❑ Minimise the use of barriers
 - ❑ Use `nowait` but only if it is safe to do so!
 - ❑ Especially minimise critical sections
 - ❑ High overhead. Can you use reduction or atomics?
 - ❑ Benchmark to find best solution
 - ❑ Experimentally try out different scheduling approaches and chunk sizes



Summary

❑ Nesting

- ❑ Operate on nested loops using OpenMP
- ❑ Compare the performance implications of different approaches for nesting

❑ Summary

- ❑ Classify permitted use of the various OpenMP clauses



- ❑ Further Reading: <https://software.intel.com/en-us/articles/32-openmp-traps-for-c-developers>