

Parallel Computing with GPUs

GPU Architectures

Part 1 – Introduction to GPUs



Dr Paul Richmond

<http://paulrichmond.shef.ac.uk/teaching/COM4521/>



This Lecture (learning objectives)

□ Introduction to GPUs

- Compare latency with throughput and identify how this relates to CPU and GPU architectures
- Identify examples from Flynn's taxonomy
- Classify the taxonomy of a GPU



Latency vs. Throughput

- ❑ Latency: The time required to perform some action
 - ❑ Measure in units of time
- ❑ Throughput: The number of actions executed per unit of time
 - ❑ Measured in units of what is produced
- ❑ E.g. *An assembly line manufactures GPUs. It takes **6 hours** to manufacture a GPU but the assembly line can manufacture **100 GPUs per day**.*



CPU vs GPU

❑ CPU



- ❑ Latency oriented
- ❑ Optimised for serial code performance
- ❑ Good for single complex tasks

❑ GPU

- ❑ Throughput oriented
- ❑ Massively parallel architecture
- ❑ Optimised for performing many similar tasks simultaneously (data parallel)



CPU vs GPU



❑ Large Cache

- ❑ Hide long latency memory access

❑ Powerful Arithmetic Logical Unit (ALU)

- ❑ Low Operation Latency

❑ Complex Control mechanisms

- ❑ Branch prediction etc.

❑ Small cache

- ❑ But faster memory throughput

❑ Energy efficient ALUs

- ❑ Long latency but high throughput

❑ Simple control


- ❑ No branch prediction

Data Parallelism

- ❑ Program has many similar threads of execution
 - ❑ Each thread performs the same behaviour on different data
 - ❑ Good for high throughput
- ❑ We can classify an architecture based on instructions and data (Flynn's Taxonomy)
 - ❑ Instructions:
 - ❑ Single instruction (SI)
 - ❑ Multiple Instruction (MI)
 - ❑ *Single Program (SP)*
 - ❑ *Multiple Program (MP)* } *Not part of the original taxonomy*
 - ❑ Data:
 - ❑ Single Data (SD) – w.r.t. *work item not necessarily single word*
 - ❑ Multiple Data (MD)
- ❑ e.g. SIMD = Single Instruction and Multiple Data




SISD and SIMD



❑ SISD

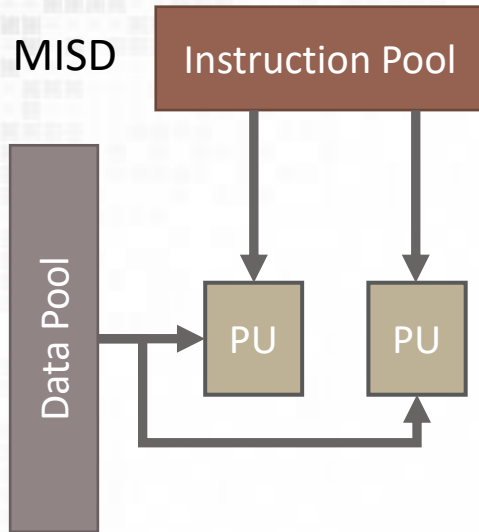
- ❑ Classic von Neumann architecture
- ❑ PU = Processing Unit



❑ SIMD


- ❑ Multiple processing elements performing the same operation simultaneously
- ❑ E.g. Early vector super computers
- ❑ Modern CPUs have SIMD instructions
 - ❑ But are not SIMD in general

MISD and MIMD



❑ MISD

❑ E.g. Pipelined architectures



❑ MIMD

❑ Processors as functionally asynchronous and independent

❑ Different processors may execute different instructions on different data

❑ E.g. Most parallel computers

❑ E.g. OpenMP programming model

SPMD and MPMD

❑ SPMD

- ❑ Multiple autonomous processors simultaneously executing a program on different data
- ❑ Program execution can have an independent path for each data point
- ❑ E.g. Message passing on distributed memory machines.

❑ MPMD

- ❑ Multiple autonomous processors simultaneously executing at least two independent programs.
- ❑ Typically client & host programming models fit this description.
- ❑ E.g. Sony PlayStation 3 SPU/PPU combination, Some system on chip configurations with CPU and GPUs





Taxonomy of a GPU

☐ What taxonomy best describes data parallelism with a GPU?

- ☐ SISD?
- ☐ SIMD?
- ☐ MISD?
- ☐ MIMD?
- ☐ SPMD?
- ☐ MPMD?



Taxonomy of a GPU

❑ What taxonomy best describes data parallelism with a GPU?

❑ **Obvious Answer: SIMD**

❑ **Less Obvious answer: SPMD**

❑ **Slightly confusing answer: SIMT (Single Instruction Multiple Thread)**

❑ This is a combination of both it differs from SIMD in that;

- 1) Each thread has its own registers
- 2) Each thread has multiple addresses
- 3) Each thread has multiple flow paths

❑ We will explore this in more detail when we look at the hardware!

❑ <http://yosefk.com/blog/simd-simt-smt-parallelism-in-nvidia-gpus.html>



Summary

❑ Introduction to GPUs

- ❑ Compare latency with throughput and identify how this relates to CPU and GPU architectures
- ❑ Identify examples from Flynn's taxonomy
- ❑ Classify the taxonomy of a GPU

❑ Next Lecture: Programming GPUs

