

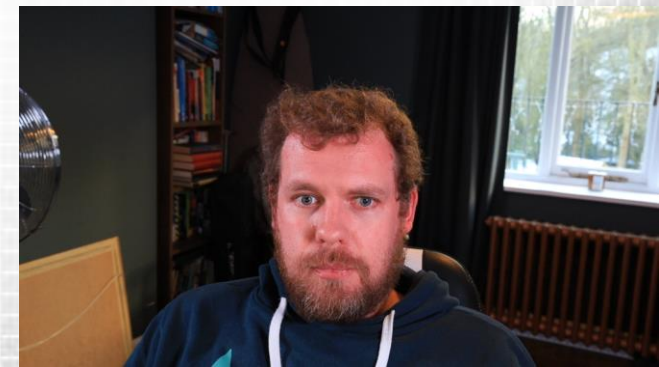
Parallel Computing with GPUs

Introduction to CUDA Part 2 – Device Code



Dr Paul Richmond

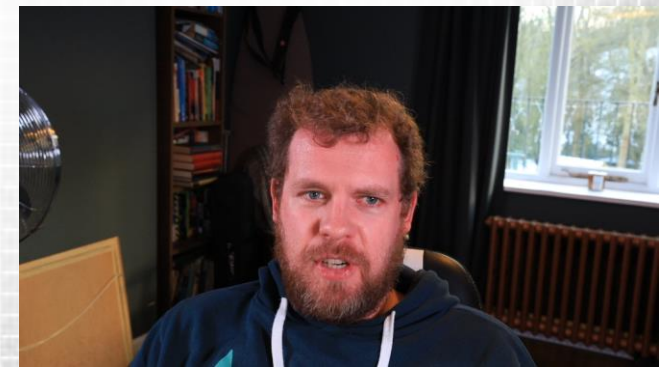
<http://paulrichmond.shef.ac.uk/teaching/COM4521/>



This Lecture (learning objectives)

❑ CUDA Device Code

- ❑ Demonstrate a simple CUDA Kernel
- ❑ Explain how the host can configure a grid of thread blocks
- ❑ Identify how the grid block configuration can be utilised by the device



A First CUDA Example

❑ Serial solution

```
for (i=0;i<N;i++) {  
    result[i] = 2*i;  
}
```

❑ We can parallelise this by assigning each iteration to a CUDA thread!



CUDA C Example: Device

```
__global__ void myKernel(int *result)
{
    int i = threadIdx.x;
    result[i] = 2*i;
}
```



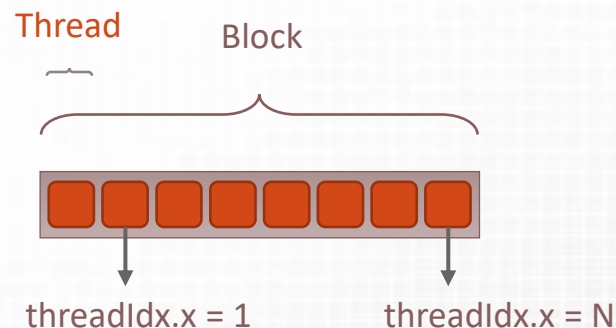
- ❑ Replace loop with a “kernel”
 - ❑ Use `__global__` specifier to indicate it is a CUDA kernel
- ❑ Use `threadIdx` dim variable to get a unique index
 - ❑ Assuming for simplicity we have only **one block** which is **1-dimensional**
 - ❑ Equivalent to your door number at CUDA Halls of Residence

CUDA C Example: Host

- ❑ Call the kernel by using the CUDA kernel launch syntax
 - ❑ `kernel<<<GRID OF BLOCKS, BLOCK OF THREADS>>>(arguments);`

```
dim3 blocksPerGrid(1,1,1);           //use only one block
dim3 threadsPerBlock(N,1,1);         //use N threads in the block
```

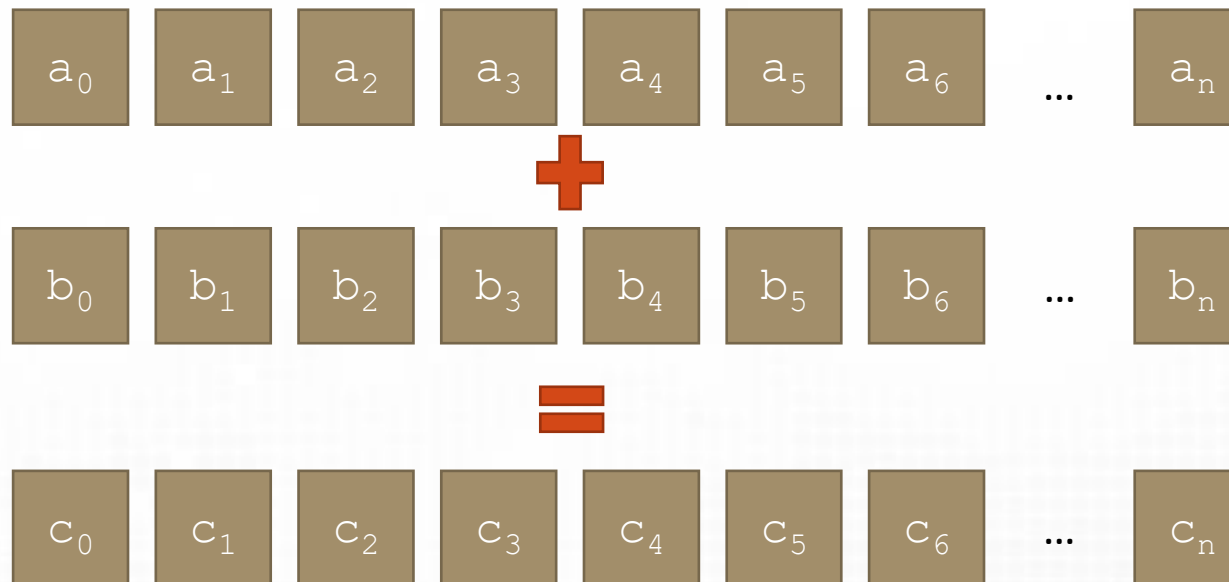
```
myKernel<<<blocksPerGrid, threadsPerBlock>>>(result);
```



Vector Addition Example

□ Consider a more interesting example

□ Vector addition: e.g. $a + b = c$



Vector Addition Example

//Kernel Code

```
__global__ void vectorAdd(float *a, float *b, float *c)
{
    int i = threadIdx.x;
    c[i] = a[i] + b[i];
}
```

//Host Code

```
...
dim3 blocksPerGrid(1,1,1);
dim3 threadsPerBlock(N,1,1); //single block of threads

vectorAdd<<<blocksPerGrid, threadsPerBlock>>>(a, b, c);
```



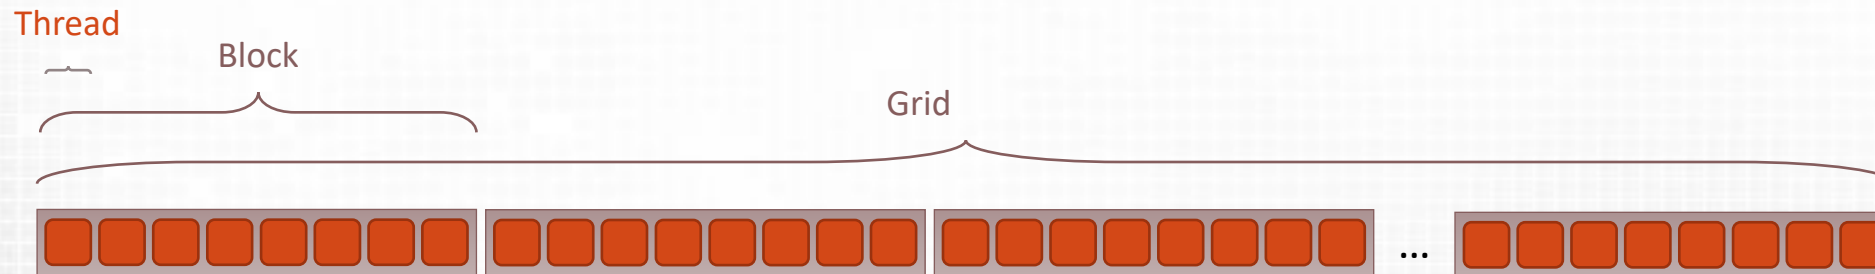
CUDA C Example: Host

- ❑ Only one block will give poor performance
 - ❑ A block gets allocated to a single SMP!
 - ❑ Solution: Use multiple blocks

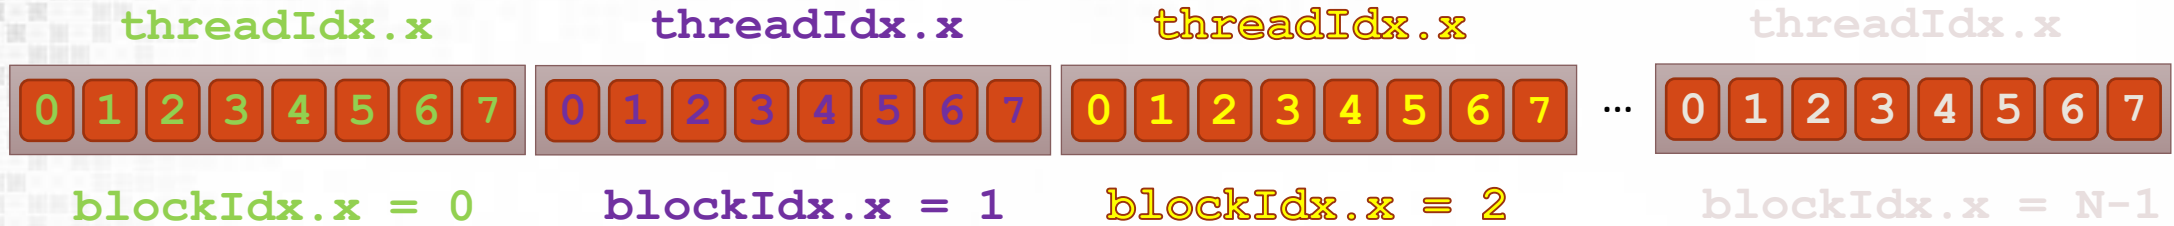


```
dim3 blocksPerGrid(N/8,1,1);    // assumes 8 divides N exactly
dim3 threadsPerBlock(8,1,1);    // 8 threads in the block
```

```
myKernel<<<blocksPerGrid, threadsPerBlock>>>(result);
```



Vector Addition Example



```
//Kernel Code
__global__ void vectorAdd(float *a, float *b, float *c)
{
    int i = blockIdx.x * blockDim.x + threadIdx.x;
    c[i] = a[i] + b[i];
}
```

- ❑ The integer `i` gives a unique thread Index used to access a unique value from the vectors `a`, `b` and `c`



A note on block sizes

- ❑ Thread block sizes can not be larger than 1024
- ❑ Max grid size is 2147483647 for 1D
 - ❑ Grid y and z dimensions are limited to 65535
- ❑ Block size should always be divisible by 32
 - ❑ This is the warp size which threads are scheduled
 - ❑ Not less than 32 as in our trivial example!
- ❑ Varying the block size will result in different performance characteristics
 - ❑ Try incrementing by values of 32 and benchmark.
- ❑ Calling a kernel with scalar parameters assumes a 1D grid of thread blocks.
 - ❑ E.g. `my_kernel<<<8, 128>>>(arguments);`



Device functions

- ❑ Kernels are always prefixed with `_global_`
- ❑ To call another function from a kernel the function must be a device function (i.e. it must be compiled for the GPU device)
 - ❑ A device function must be prefixed with `_device_`
- ❑ A device function is not available from the host
 - ❑ Unless it is also prefixed with `_host_`



```
int increment(int a){ return a + 1; }
```

Host only

```
__device__ int increment(int a){ return a + 1; }
```

Device only

```
__device__ __host__ int increment(int a){ return a + 1; }
```

Host and device

Global functions are always `void` return type

Summary

❑ CUDA Device Code

- ❑ Demonstrate a simple CUDA Kernel
- ❑ Explain how the host can configure a grid of thread blocks
- ❑ Identify how the grid block configuration can be utilised by the device

❑ Next Lecture: Host Code and Memory Management



Acknowledgements and Further Reading

❑ Some of the content in this lecture material has been provided by;

1. GPUComputing@Sheffield Introduction to CUDA Teaching Material

❑ Originally from content provided by Alan Gray at EPCC/NVIDIA

2. NVIDIA Educational Material

❑ Specifically Mark Harris's (Introduction to CUDA C)

❑ Further Reading

❑ Essential Reading: CUDA C Programming Guide

❑ <http://docs.nvidia.com/cuda/cuda-c-programming-guide/>