


Parallel Computing with GPUs

CUDA Memory Part 3 – Read Only and Texture Memory



Dr Paul Richmond

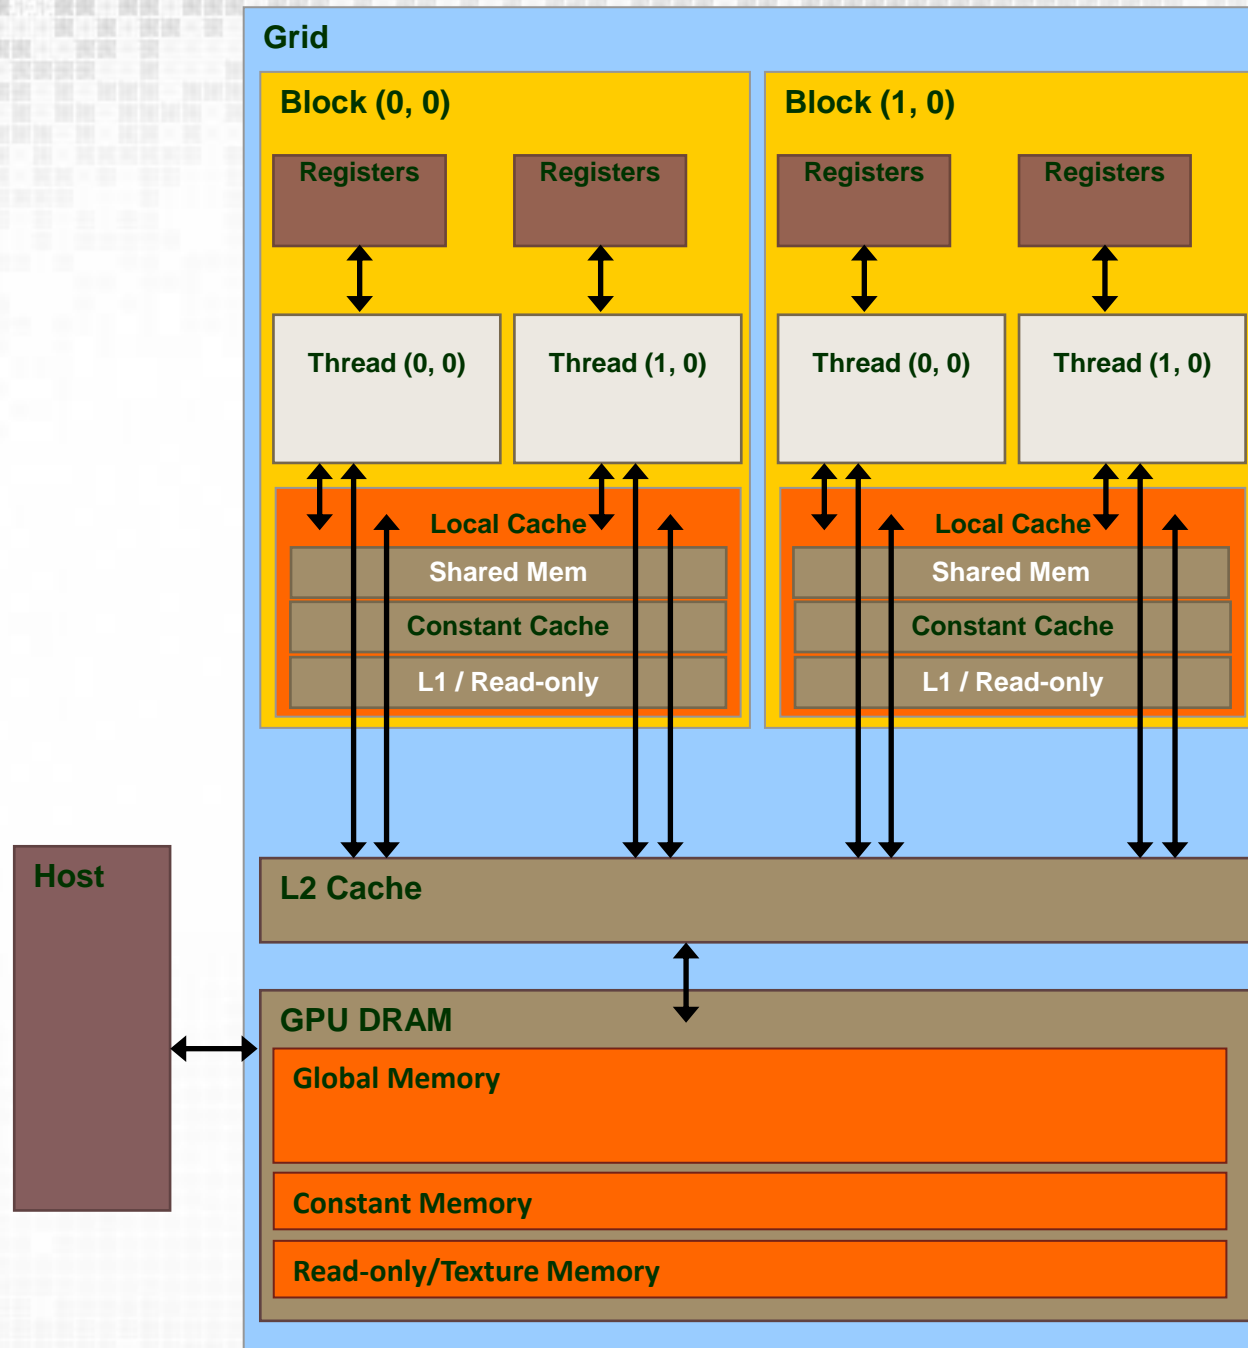
<http://paulrichmond.shef.ac.uk/teaching/COM4521/>



This Lecture (learning objectives)

- ❑ Read Only and Texture Memory
 - ❑ Identify use cases for read only and texture memory
 - ❑ Demonstrate texture memory binding
 - ❑ Highlight the simplicity of read on memory usage
 - ❑ Extra Material: Demonstrate Bindless Textures





Read-only and Texture Memory



- ❑ Separate in Kepler but unified with L1 thereafter
 - ❑ Same use case but used in different ways

- ❑ When to use read-only or texture
 - ❑ When data is read only
 - ❑ Good for bandwidth limited kernels
 - ❑ Regular memory accesses with good locality (think about the way textures are accessed)
 - ❑ Texture cache can outperform read only cache for certain scenarios
 - ❑ Normalisation/interpolation
 - ❑ 2D and 3D loads
 - ❑ Read only cache can outperform texture cache
 - ❑ Loads of 4 byte values

- ❑ Two Methods for utilising Read-only/Texture Memory
 - ❑ Bind memory to texture (or use advanced bindless textures in CUDA 5.0+)
 - ❑ Hint the compiler to load via read-only cache


Texture Memory Binding

□ Known as bound texture (or texture reference method)

```
#define N 1024
texture<float, 1, cudaReadModeElementType> tex;

__global__ void kernel() {
    int i = blockIdx.x * blockDim.x + threadIdx.x;
    float x = tex1Dfetch(tex, i);
}

int main() {
    float *buffer;
    cudaMalloc(&buffer, N*sizeof(float));
    cudaBindTexture(0, tex, buffer, N*sizeof(float));
    kernel << <grid, block >> >();
    cudaUnbindTexture(tex);
    cudaFree(buffer);
}
```



Texture Memory Binding

❑ Known as bound texture (or texture reference method)

```
#define N 1024
texture<float> 1, cudaReadModeElementType> tex;

__global__ void kernel() {
    int i = blockIdx.x * blockDim.x + threadIdx.x;
    float x = tex1Dfetch(tex, i);
}


int main() {
    float *buffer;
    cudaMalloc(&buffer, N*sizeof(float));
    cudaBindTexture(0, tex, buffer, N*sizeof(float));
    kernel << <grid, block >> >();
    cudaUnbindTexture(tex);
    cudaFree(buffer);
}
```

Must be either;

❑ char, short, long,
long long, float or
double

Vector Equivalents are also
permitted e.g.

❑ uchar4



Texture Memory Binding

❑ Known as bound texture (or texture reference method)

```
#define N 1024
texture<float, 1, cudaReadModeElementType> tex;

__global__ void kernel() {
    int i = blockIdx.x * blockDim.x + threadIdx.x;
    float x = tex1Dfetch(tex, i);
}

int main() {
    float *buffer;
    cudaMalloc(&buffer, N*sizeof(float));
    cudaBindTexture(0, tex, buffer, N*sizeof(float));
    kernel << <grid, block >> >();
    cudaUnbindTexture(tex);
    cudaFree(buffer);
}
```

Dimensionality:

- ❑ cudaTextureType1D (1)
- ❑ cudaTextureType2D (2)
- ❑ cudaTextureType3D (3)
- ❑ cudaTextureType1DLayered (4)
- ❑ cudaTextureType2DLayered (5)
- ❑ cudaTextureTypeCubemap (6)
- ❑ cudaTextureTypeCubemapLayered (7)



Texture Memory Binding

❑ Known as bound texture (or texture reference method)

```
#define N 1024
texture<float, 1, cudaReadModeElementType> tex;

__global__ void kernel() {
    int i = blockIdx.x * blockDim.x + threadIdx.x;
    float x = tex1Dfetch(tex, i);
}

int main() {
    float *buffer;
    cudaMalloc(&buffer, N*sizeof(float));
    cudaBindTexture(0, tex, buffer, N*sizeof(float));
    kernel << <grid, block >> >();
    cudaUnbindTexture(tex);
    cudaFree(buffer);
}
```

Value normalization:

- ❑ cudaReadModeElementType
- ❑ cudaReadModeNormalizedFloat
 - ❑ Normalises values across range



Texture Memory Binding on 2D Arrays

```
#define N 1024
texture<float, 2, cudaReadModeElementType> tex;

__global__ void kernel() {
    int x = blockIdx.x * blockDim.x + threadIdx.x;
    int y = blockIdx.y * blockDim.y + threadIdx.y;
    float v = tex2D (tex, x, y);
}

int main() {
    float *buffer;
    cudaMalloc(&buffer, W*H*sizeof(float));
    cudaChannelFormatDesc desc = cudaCreateChannelDesc<float>();
    cudaBindTexture2D(0, tex, buffer, desc, W,
                      H, W*sizeof(float));
    kernel << <grid, block >> >();
    cudaUnbindTexture(tex);
    cudaFree(buffer);
}
```

❑ Use tex2D rather than tex1Dfetch for CUDA arrays

❑ Note that last arg of **cudaBindTexture2D** is pitch

❑ Row size not != total size



Read-only Memory

- ❑ No textures required
- ❑ Hint to the compiler that the data is read-only without pointer aliasing
 - ❑ Using the `const` and `__restrict__` qualifiers
 - ❑ Suggests the compiler should use `__ldg` but does not guarantee it
- ❑ Not the same as `__constant__`
 - ❑ Does not require broadcast reading

```
#define N 1024

__global__ void kernel(float const* __restrict__ buffer) {
    int i = blockIdx.x * blockDim.x + threadIdx.x;
    float x1 = buffer[i];
    float x2 = __ldg(buffer[i]);
}

int main() {
    float *buffer;
    cudaMalloc(&buffer, N*sizeof(float));
    kernel << <grid, block >> >(buffer);
    cudaFree(buffer);
}
```

Probably read through read only cache
Definitely read through read only cache



Summary

- ❑ Read Only and Texture Memory
 - ❑ Identify use cases for read only and texture memory
 - ❑ Demonstrate texture memory binding
 - ❑ Highlight the simplicity of read on memory usage
 - ❑ Extra Material: Demonstrate Bindless Textures



Acknowledgements and Further Reading

- ❑ <http://devblogs.nvidia.com/paralleforall/cuda-pro-tip-kepler-texture-objects-improve-performance-and-flexibility/>
- ❑ Mike Giles (Oxford): Different Memory and Variable Types
 - ❑ <https://people.maths.ox.ac.uk/gilesm/cuda/>
- ❑ CUDA Programming Guide
 - ❑ <http://docs.nvidia.com/cuda/cuda-c-programming-guide/#texture-memory>



Bindless Textures (Advanced)

```
#define N 1024

__global__ void kernel(cudaTextureObject_t tex) {
    int i = blockIdx.x * blockDim.x + threadIdx.x;
    float x = tex1Dfetch(tex, i);
}

int main() {
    float *buffer;
    cudaMalloc(&buffer, N*sizeof(float));

    cudaResourceDesc resDesc;
    memset(&resDesc, 0, sizeof(resDesc));
    resDesc.resType = cudaResourceTypeLinear;
    resDesc.res.linear.devPtr = buffer;
    resDesc.res.linear.desc.f = cudaChannelFormatKindFloat;
    resDesc.res.linear.desc.x = 32; // bits per channel
    resDesc.res.linear.sizeInBytes = N*sizeof(float);

    cudaTextureDesc texDesc;
    memset(&texDesc, 0, sizeof(texDesc));
    texDesc.readMode = cudaReadModeElementType;

    cudaTextureObject_t tex;
    cudaCreateTextureObject(&tex, &resDesc, &texDesc, NULL);
    kernel << <grid, block >> >(tex);
    cudaDestroyTextureObject(tex);
    cudaFree(buffer);
}
```

- ❑ Texture Object Approach (Kepler+ and CUDA 5.0+)
- ❑ Textures only need to be created once
 - ❑ No need for binding an unbinding
- ❑ Better performance than binding
 - ❑ Small kernel overhead
- ❑ More details in programming guide
 - ❑ <http://docs.nvidia.com/cuda/cuda-c-programming-guide/index.html#texture-object-api>

Address and Filter Modes (Bindless Textures)

- ❑ `addressMode`: Dictates what happens when addresses are out of bounds. E.g.
 - ❑ `cudaAddressModeClamp`: in which case addresses out of bounds will be clamped to range
 - ❑ `cudaAddressModeWrap`: in which case addresses out of bounds will wrap
- ❑ `filterMode`: Allows values read from the texture to be filtered. E.g.
 - ❑ `cudaFilterModeLinear`: Linearly interpolates between points
 - ❑ `cudaFilterModePoint`: Gives the value at the specific texture point

```
cudaTextureObject_t tex;  
cudaCreateTextureObject(&tex, &resDesc, &texDesc, NULL);  
tex.addressMode = cudaAddressModeClamp;
```

Bindless Textures

```
texture<float, 1, cudaReadModeElementType> tex;  
tex.addressMode = cudaAddressModeClamp;
```

Bound Textures