

Parallel Computing with GPUs

Shared Memory Part 1 – Introduction to Shared Memory



The
University
Of
Sheffield.

Dr Paul Richmond

<http://paulrichmond.shef.ac.uk/teaching/COM4521/>



This Lecture (learning objectives)

❑ Shared Memory

- ❑ Repeat important concepts of grids, blocks and warps
- ❑ Identify a use case for shared memory
- ❑ Demonstrate the use of shared memory on a simple problem
- ❑ Recognise potential issues caused by use of shared memory (boundary conditions)

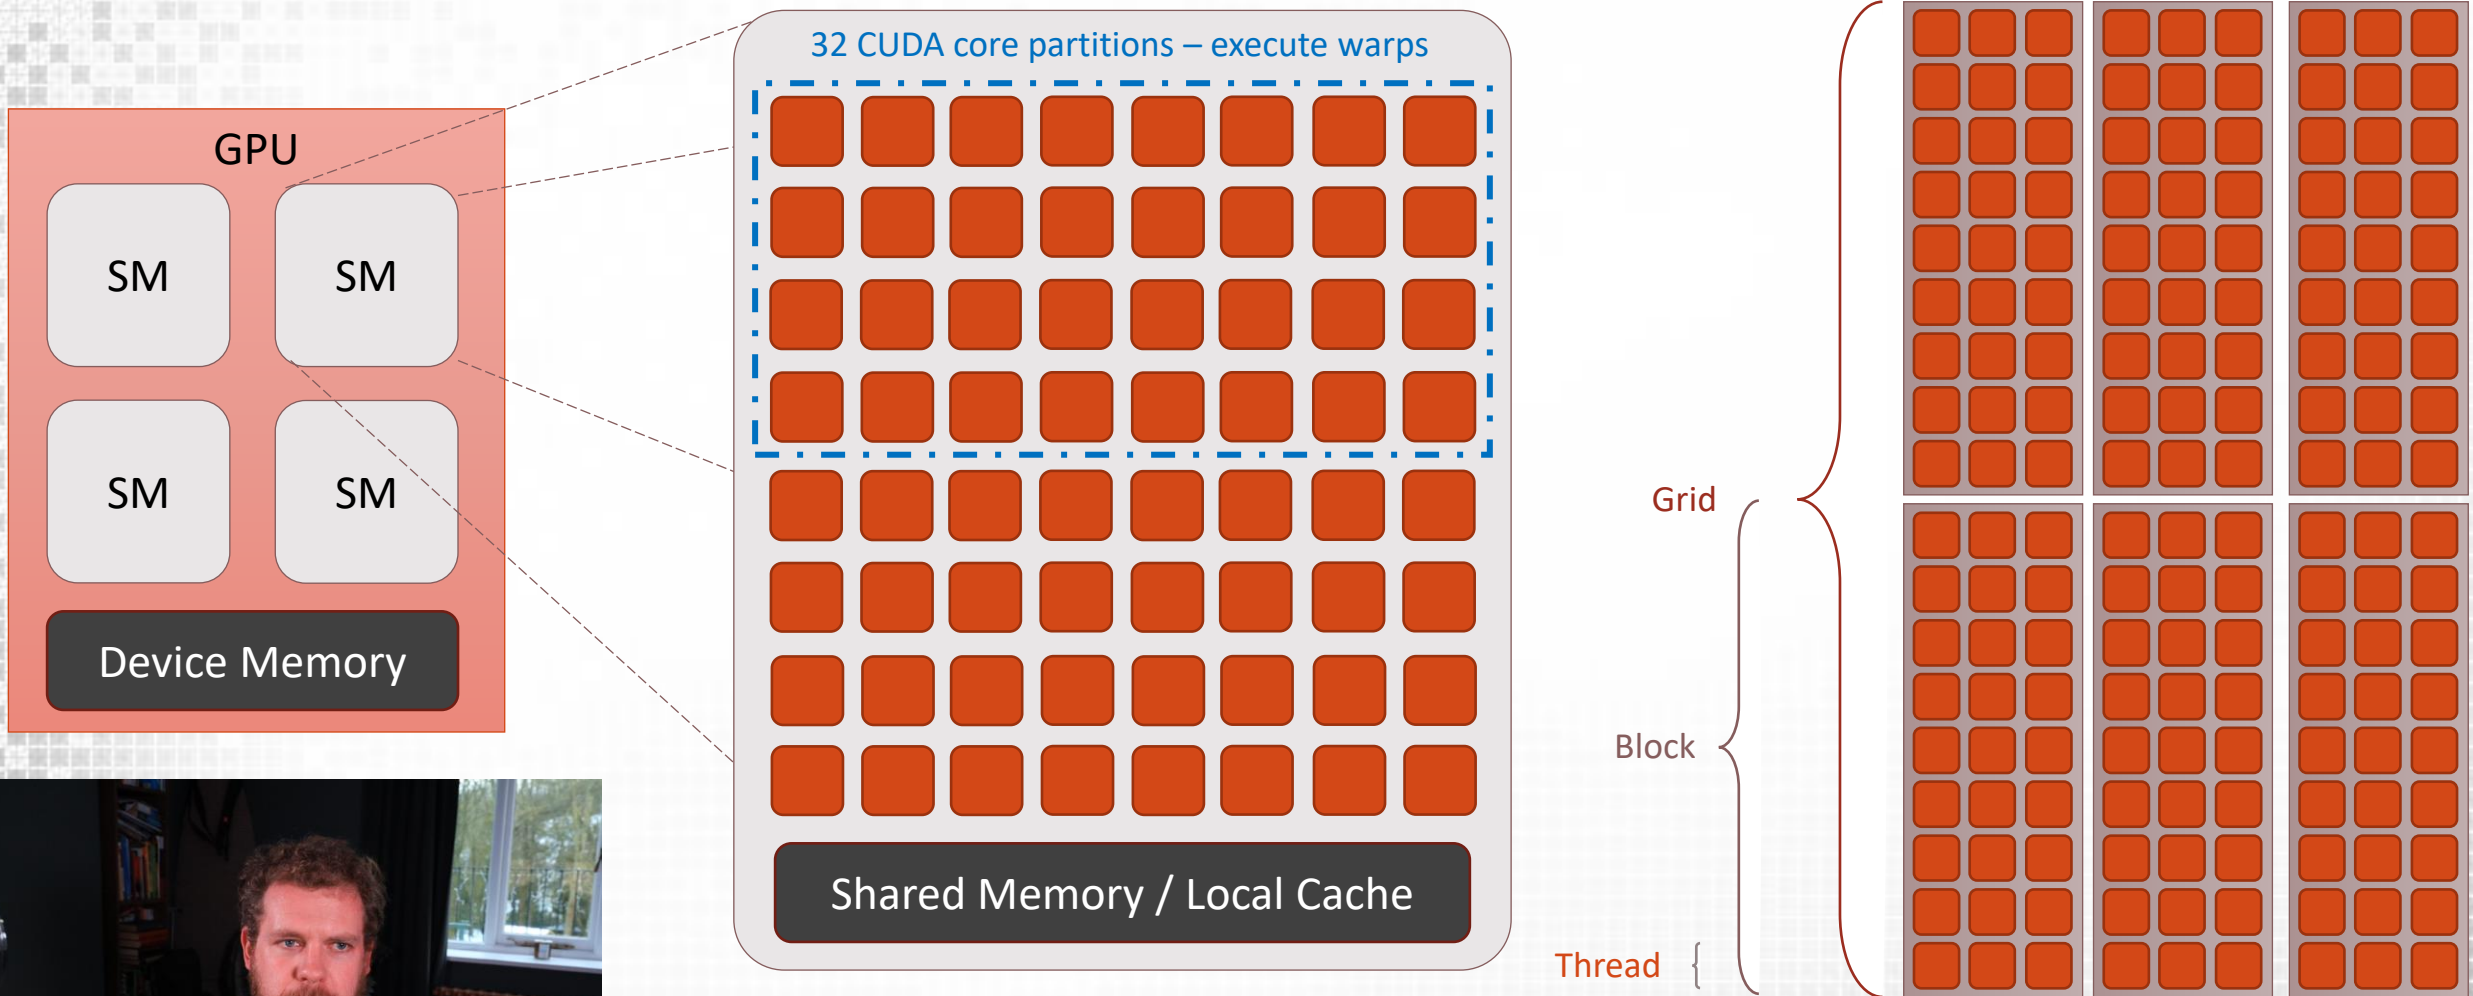


Review of last week

- ❑ We have seen the importance of different types of memory
 - ❑ And observed the performance improvement from read-only and constant cache usage
- ❑ So far we have seen how CUDA can be used for performing thread local computations; e.g.
 - ❑ Load data from memory to registers
 - ❑ Perform thread-local computations
 - ❑ Store results back to global memory
- ❑ We will now consider another important type of memory
 - ❑ Shared memory



Grids, Blocks, Warps & Threads



Shared Memory

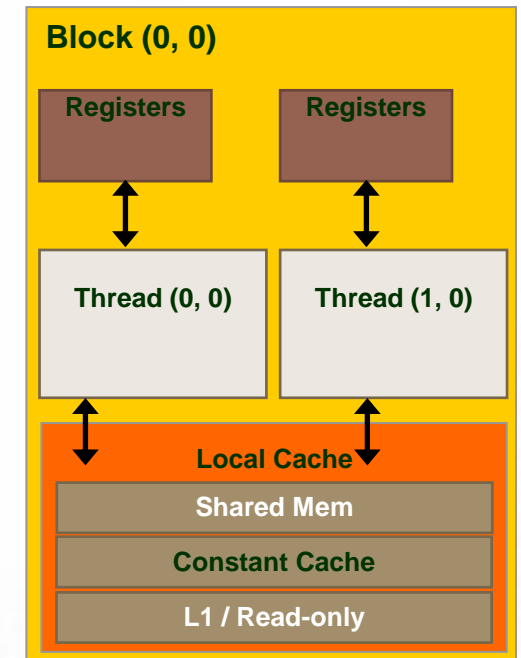
❑ Architecture Details

- ❑ In Maxwell 64KB of Shared Memory is dedicated

❑ Its just another Cache, right?

- ❑ User configurable
- ❑ Requires manually loading and synchronising data

Maxwell/Pascal



Shared Memory

❑ Performance

- ❑ Shared memory is very fast
- ❑ Bandwidth > 1 TB/s


❑ Block level computation

- ❑ Challenges the thread level view...
- ❑ Allows data to be shared between threads in the same block
- ❑ User configurable cache at the thread block level
- ❑ Still no broader synchronisation beyond the level of thread blocks



Block Local Computation

- ❑ Partition data into groups that fit into shared memory
- ❑ Load subset of data into shared memory
- ❑ Perform computation on the subset
- ❑ Copy subset back to global memory



Move, execute, move

Thread level parallelism

- ☐ From Host view
 - ☐ Move: Data to GPU memory
 - ☐ Execute: Kernel
 - ☐ Move: Data back to host
- ☐ From Device view
 - ☐ Move: Data from device memory to registers
 - ☐ Execute: instructions
 - ☐ Move: Data back to device memory

Block level parallelism

- ☐ From Host view
 - ☐ Move: Data to GPU memory
 - ☐ Execute: Kernel
 - ☐ Move: Data back to host
- ☐ From Device view
 - ☐ Move: Data from device memory to local cache
 - ☐ Execute: subset of kernel (reusing cached values)
 - ☐ Move: Data back to device memory
- ☐ From Block View
 - ☐ Move: Data from local cache
 - ☐ Execute: instructions
 - ☐ Move: Data back to local cache (or device memory)



A Case for Shared Memory

```
__global__ void sum3_kernel(int *c, int *a)
{
    int i = blockIdx.x*blockDim.x + threadIdx.x;
    int left, right;

    //load value at i-1
    left = 0;
    if (i > 0)
        left = a[i - 1];

    //load value at i+1
    right = 0;
    if (i < (N - 1))
        right = a[i + 1];

    c[i] = left + a[i] + right; //sum three values
}
```

Do we have a candidate for block level parallelism using shared memory?



A Case for Shared Memory

```
__global__ void sum3_kernel(int *c, int *a)
{
    int i = blockIdx.x*blockDim.x + threadIdx.x;
    int left, right;

    //load value at i-1
    left = 0;
    if (i > 0)
        left = a[i - 1];

    //load value at i+1
    right = 0;
    if (i < (N - 1))
        right = a[i + 1];

    c[i] = left + a[i] + right; //sum three values
}
```

- ❑ Currently: Thread-local computation
- ❑ Bandwidth limited
 - ❑ Requires three loads per thread (at index $i-1$, i , and $i+1$)
- ❑ Block level solution: load each value only once!



CUDA Shared memory

- ❑ Shared memory between threads in the same block can be defined using `__shared__`
- ❑ Shared variables are only accessible from within device functions
 - ❑ Not addressable in host code
- ❑ Must be careful to avoid race conditions
 - ❑ Multiple threads writing to the same shared memory variable
 - ❑ Results in undefined behaviour
 - ❑ Typically write to shared memory using `threadIdx`
 - ❑ Thread level synchronisation is available through `__syncthreads()`
 - ❑ Synchronises threads in the block

```
__shared__ int s_data[BLOCK_SIZE];
```



Example

```
__global__ void sum3_kernel(int *c, int *a)
{
    __shared__ int s_data[BLOCK_SIZE];

    int i = blockIdx.x*blockDim.x + threadIdx.x;
    int left, right;

    s_data[threadIdx.x] = a[i];
    __syncthreads();

    //load value at i-1
    left = 0;
    if (i > 0){
        left = s_data[threadIdx.x - 1];
    }

    //load value at i+1
    right = 0;
    if (i < (N - 1)){
        right = s_data[threadIdx.x + 1];
    }

    c[i] = left + s_data[threadIdx.x] + right; //sum
}
```

What is wrong with this code?

- ❑ Allocate a shared array
 - ❑ One integer element per thread
- ❑ Each thread loads a single item to shared memory
- ❑ Call `__syncthreads` to ensure shared memory data is populated by all threads
- ❑ Load all elements through shared memory



Example

```
__global__ void sum3_kernel(int *c, int *a)
{
    __shared__ int s_data[BLOCK_SIZE];

    int i = blockIdx.x*blockDim.x + threadIdx.x;
    int left, right;

    s_data[threadIdx.x] = a[i];
    __syncthreads();

    //load value at i-1
    left = 0;
    if (i > 0){
        if (threadIdx.x > 0)
            left = s_data[threadIdx.x - 1];
        else
            left = a[i - 1];
    }

    //load value at i+1
    right = 0;
    if (i < (N - 1)){
        if (threadIdx.x < (BLOCK_SIZE-1))
            right = s_data[threadIdx.x + 1];
        else
            right = a[i + 1];
    }

    c[i] = left + s_data[threadIdx.x] + right; //sum
}
```

❑ Additional step required!

❑ Check boundary conditions for the edge of the block



Problems with Shared memory

- ❑ In the example we saw the introduction of boundary conditions
 - ❑ Global loads still present at boundaries
 - ❑ We have introduced divergence in the code (remember the SIMD model)
 - ❑ This is even more prevalent in 2D examples where we *tile* data into shared memory

```
//boundary condition
left = 0;
if (i > 0){
    if (threadIdx.x > 0)
        left = s_data[threadIdx.x - 1];
    else
        left = a[i - 1];
}
```



Dynamically Assigned Shared Memory

- ❑ It is possible to dynamically assign shared memory at runtime.
- ❑ Requires both a host and device modification to code
 - ❑ Device: Must declare shared memory as extern
 - ❑ Host: Must declare shared memory size in kernel launch parameters

```
unsigned int sm_size = sizeof(float)*DIM*DIM;  
image_kernel<<<blocksPerGrid, threadsPerBlock, sm_size >>>(d_image);  
  
__global__ void image_kernel(float *image)  
{  
    extern __shared__ float s_data[];  
}
```

Is equivalent to

```
image_kernel<<<blocksPerGrid, threadsPerBlock>>>(d_image);  
  
__global__ void image_kernel(float *image)  
{  
    __shared__ float *s_data[DIM][DIM];  
}
```



Summary

❑ Shared Memory

- ❑ Repeat important concepts of grids, blocks and warps
- ❑ Identify a use case for shared memory
- ❑ Demonstrate the use of shared memory on a simple problem
- ❑ Recognise potential issues caused by use of shared memory (boundary conditions)

❑ Next Lecture: Shared Memory Bank Conflicts

