

Parallel Computing with GPUs

Shared Memory Part 2 – Bank Conflicts



Dr Paul Richmond

<http://paulrichmond.shef.ac.uk/teaching/COM4521/>



This Lecture (learning objectives)

❑ Shared Memory Bank Conflicts

- ❑ Explain the concepts of memory banks and how conflicts can occur
- ❑ Give examples of access strides and the impact on bank conflicts
- ❑ Present solutions for avoiding bank conflicts




Shared Memory Bank Conflicts

- ❑ Shared memory is arranged into 4byte (32bit banks)
 - ❑ A load or store of N addresses spanning N distinct banks can be serviced simultaneously
 - ❑ Overall bandwidth of $\times N$ a single module
 - ❑ Kepler+ can also serve broadcast accesses simultaneously
- ❑ A bank conflict occurs when two threads request addresses from the same bank (without reading in broadcast pattern)
 - ❑ Results in serialisation of the access
- ❑ Bank conflicts only occur between threads in a warp
 - ❑ There are 32 banks and 32 threads per warp
 - ❑ If two threads in a warp access the same bank this is said to be a 2-way bank conflict



Think about you block sized array of floats
bank = (index * stride) % 32

Logical View of Shared Memory Banks



Access Strides

❑ Stride refers to the size (in increments of the bank size) between each threads memory access pattern

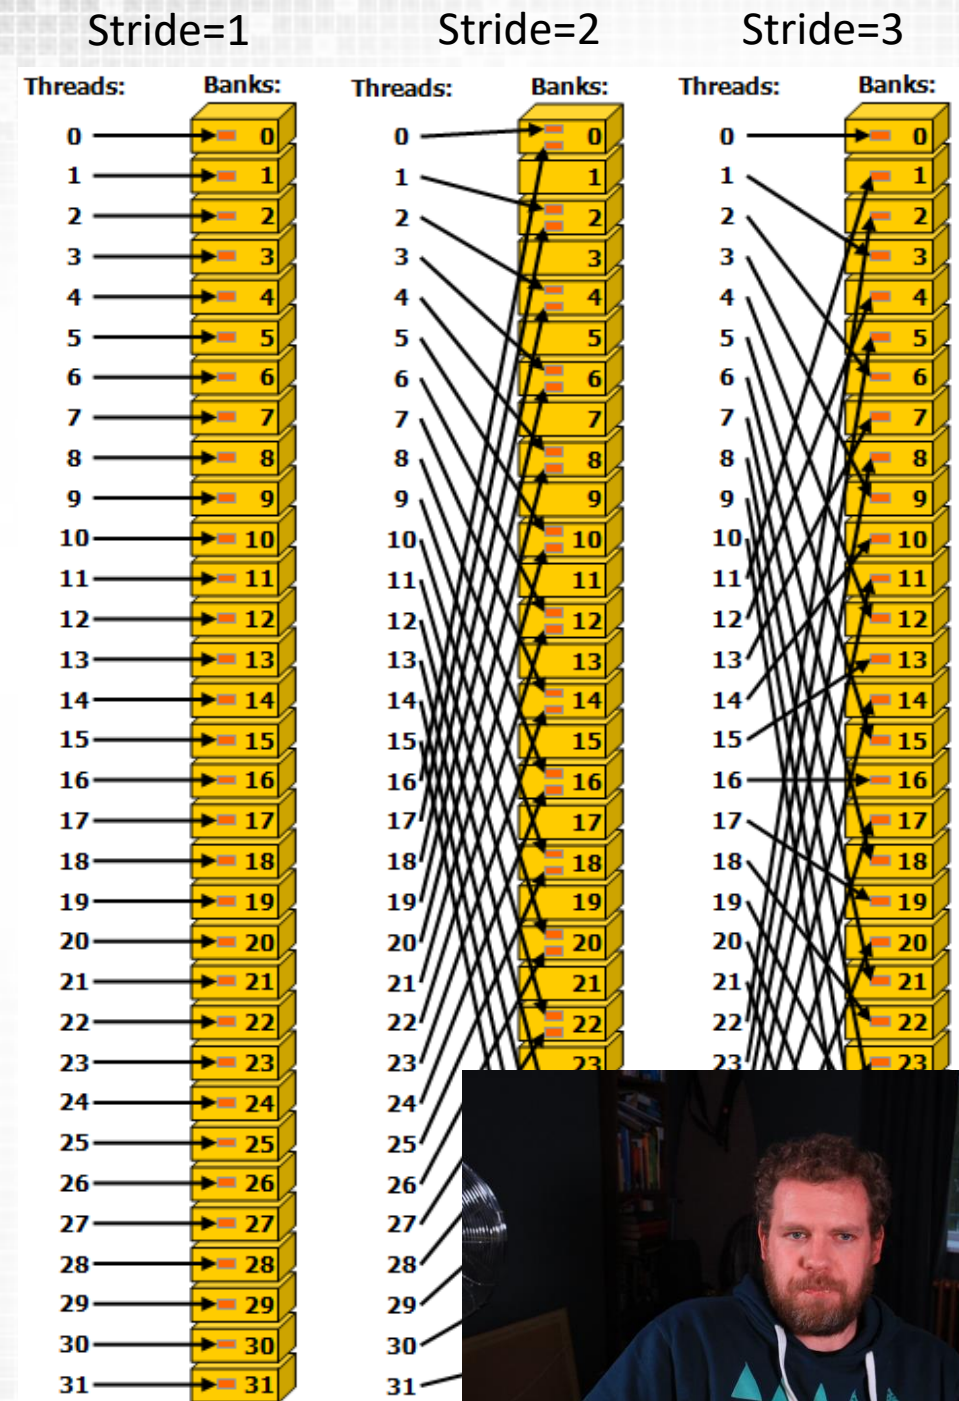
❑ If threads access consecutive 4 byte values (e.g. `int` or `float`) then the stride is 1.

❑ No conflicts

❑ If a threads accesses consecutive 8 bytes values (e.g. `double`) then the stride is 2.


❑ 2 way conflicts

❑ In general odd strides result in no conflicts



Stride (4 byte)	1		
TPB	128		
threadIdx.x		index	bank
0		0	1
1		1	2
2		2	3
3		3	4
4		4	5
5		5	6
6		6	7
7		7	8
8		8	9
9		9	10
10		10	11
31		31	12
		Banks Used	32
		Max Conflicts	1

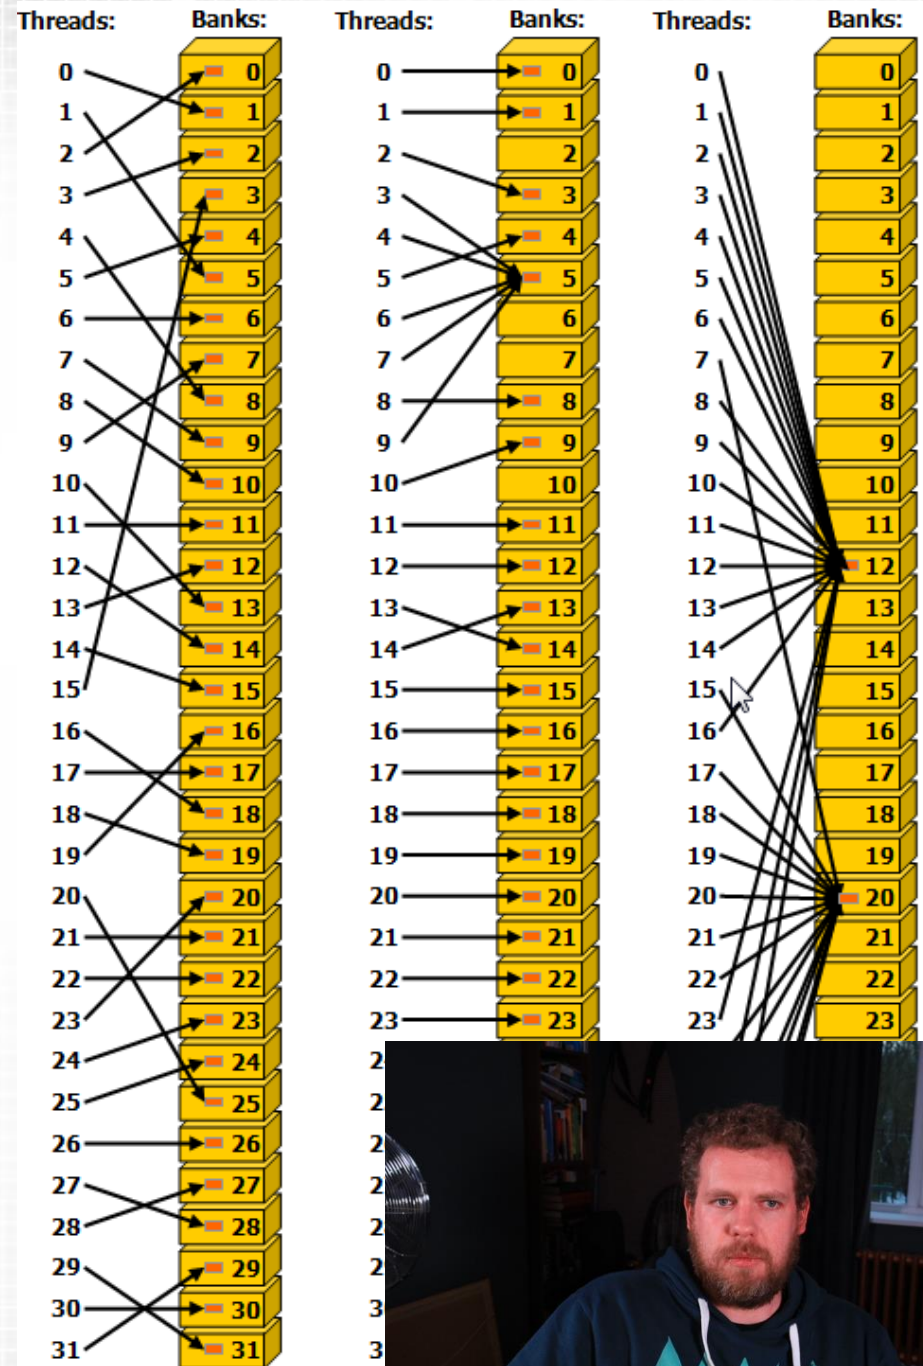
$\text{bank} = (\text{index} * \text{stride}) \% 32$



More on SM banks

- ❑ Irregular access is fine as long as no bank conflicts
- ❑ Multiple threads can access the same bank conflict free **if** they access the same addresses (e.g. a broadcast access pattern)
- ❑ Broadcast can be to any number of threads in a warp

```
__shared__ float s_data[??];  
//read from shared memory using broadcast  
  
some_thread_value = s_data[0] ;
```



Strided access example

```
__shared__ float s_data[BLOCK_SIZE];  
  
//load or calculate some_thread_value  
  
s_data[threadIdx.x*2] = some_thread_value;  
__syncthreads();
```


Thread	Bank
0	0
1	1
2	2
3	3
4	4
5	5
6	6
7	7
...	...
31	31

- ❑ What is the stride?
- ❑ What is the level of conflict?
- ❑ How can this be improved?



Strided access example

```
__shared__ float s_data[BLOCK_SIZE*2];  
  
//load or calculate some_thread_value  
  
s_data[threadIdx.x*2] = some_thread_value;  
__syncthreads();
```




- ❑ What is the stride? **2**
- ❑ What is the level of conflict? **2 way**
- ❑ How can this be improved? **Remove the stride (use a stride of 1)**



Edge case for broadcast access

```
__shared__ char s_data[BLOCK_SIZE];  
  
//load or calculate some_thread_value  
  
s_data[threadIdx.x] = some_thread_value;  
__syncthreads();
```




Edge case for broadcast access

```
__shared__ char s_data[BLOCK_SIZE];

//load or calculate some_thread_value

s_data[threadIdx.x] = some_thread_value;
__syncthreads();
```



❑ Stride: 0.25 BUT no conflict (for compute capability > 2.0)

❑ Why?

❑ “A shared memory request for a warp does not generate a bank conflict between two threads that access any sub-word within the same 32-bit word or within two 32-bit words whose indices i and j are in the same 64-word aligned segment (i.e., a segment whose first index is a multiple of 64) and such that $j=i+32$ (even though the addresses of the two sub-words fall in the same bank): In that case, for read accesses, the 32-bit words are broadcast to the requesting threads and for write accesses, each sub-word is written by only one of the threads (which thread performs the write is undefined)”. - <https://docs.nvidia.com/cuda/cuda-c-programming-guide/index.html#shared-memory-3-0>

❑ E.g. Adjacent threads accessing the same bank falls under the broadcast pattern



Large Data Types and Phasing

- ❑ For larger data types Volta+ Hardware allows scheduling of address accesses in phases
 - ❑ Bank conflicts only occur between threads in the same phase
- ❑ Double and 8B data types
 - ❑ 2 Phases
 - ❑ Process addresses in first 16 threads of warp
 - ❑ Process addresses in the last 16 of the warp
- ❑ 16B data types (E.g. float4)
 - ❑ 4 Phases
 - ❑ Each phase processes addresses in a quarter of the warp



Bank conflicts with 2D tiles

```
__global__ void image_kernel(float *image)
{
    __shared__ float s_data[BLOCK_DIM][BLOCK_DIM];

    for (int i = 0; i < BLOCK_DIM; i++) {
        some_thread_value += f(s_data[threadIdx.x][i]);
    }
}
```

bank = threadIdx.x * stride % 32

- ❑ Example where each thread (of 2D block) operates on a row
 - ❑ Loads values by column

Shared Memory Bank

0	1	2	3	4	5	6	7		31
0	1	2	3	4	5	6	7		31
0	1	2	3	4	5	6	7		31
0	1	2	3	4	5	6	7		31
0	1	2	3	4	5	6	7		31
0	1	2	3	4	5	6	7		31
0	1	2	3	4	5	6	7		31
0	1	2	3	4	5	6	7		31
0	1	2	3	4	5	6	7		0



Bank conflicts with 2D tiles

```
__global__ void image_kernel(float *image)
{
    __shared__ float s_data[BLOCK_DIM][BLOCK_DIM];

    for (int i = 0; i < BLOCK_DIM; i++) {
        some_thread_value += f(s_data[threadIdx.x][i]);
    }
}
```

bank = threadIdx.x * stride % 32

BLOCK_DIM=32

i=0

Shared Memory Bank

0	1	2	3	4	5	6	7		31
0	1	2	3	4	5	6	7		31
0	1	2	3	4	5	6	7		31
0	1	2	3	4	5	6	7		31
0	1	2	3	4	5	6	7		31
0	1	2	3	4	5	6	7		31
0	1	2	3	4	5	6	7		31
0	1	2	3	4	5	6	7		31
0	1	2	3	4	5	6	7		0

- ❑ Example where each thread (of 2D block) operates on a row
 - ❑ Loads values by column

- ❑ 32 way bank conflicts!
 - ❑ Very bad
- ❑ Stride = 32



Bank conflicts with 2D tiles

```
__global__ void image_kernel(float *image)
{
    __shared__ float s_data[BLOCK_DIM][BLOCK_DIM];

    for (int i = 0; i < BLOCK_DIM; i++) {
        some_thread_value += f(s_data[threadIdx.x][i]);
    }
}
```

bank = threadIdx.x * stride % 32

BLOCK_DIM=32

i=0

Shared Memory Bank

0	1	2	3	4	5	6	7		31
0	1	2	3	4	5	6	7		31
0	1	2	3	4	5	6	7		31
0	1	2	3	4	5	6	7		31
0	1	2	3	4	5	6	7		31
0	1	2	3	4	5	6	7		31
0	1	2	3	4	5	6	7		31
0	1	2	3	4	5	6	7		31
0	1	2	3	4	5	6	7		0

- ❑ Example where each thread (of 2D block) operates on a row
 - ❑ Loads values by column

- ❑ How to fix
 - ❑ Memory padding
 - ❑ Transpose the matrix
 - ❑ Or operate on columns (loading by row) if possible



Bank conflicts with 2D tiles

```
__global__ void image_kernel(float *image)
{
    __shared__ float s_data[BLOCK_DIM][BLOCK_DIM+1];

    for (int i = 0; i < BLOCK_DIM; i++){
        some_thread_value += f(d_data[threadIdx.x][i]);
    }
}
```

bank = threadIdx.x * stride % 32

BLOCK_DIM+1=33

Shared Memory Bank

0	1	2	3	4	5	6	7		0
1	2	3	4	5	6	7	8		1
2	3	4	5	6	7	8	9		2
3	4	5	6	7	8	9	10		3
4	5	6	7	8	9	10	11		4
5	6	7	8	9	10	11	12		5
6	7	8	9	10	11	12	13		6
7	8	9	10	11	12	13	14		7
31	0	1	2	3	4	5	6		31

Memory Padding Solution



Bank conflicts with 2D tiles

```
__global__ void image_kernel(float *image)
{
    __shared__ float s_data[BLOCK_DIM][BLOCK_DIM+1];

    for (int i = 0; i < BLOCK_DIM; i++){
        some_thread_value += f(d_data[threadIdx.x][i]);
    }
}
```

bank = threadIdx.x * stride % 32

BLOCK_DIM+1=33

i=0

Shared Memory Bank

0	1	2	3	4	5	6	7		0
1	2	3	4	5	6	7	8		1
2	3	4	5	6	7	8	9		2
3	4	5	6	7	8	9	10		3
4	5	6	7	8	9	10	11		4
5	6	7	8	9	10	11	12		5
6	7	8	9	10	11	12	13		6
7	8	9	10	11	12	13	14		7
31	0	1	2	3	4	5	6		31

❑ Memory Padding Solution



❑ Every thread in warp reads from different bank

❑ *Alternative: Transpose solution left to you!*

Summary

❑ Shared Memory Bank Conflicts

- ❑ Explain the concepts of memory banks and how conflicts can occur
- ❑ Give examples of access strides and the impact on bank conflicts
- ❑ Present solutions for avoiding bank conflicts

❑ Next Lecture: Boundary Conditions



Kepler Shared Memory Bank Size

- ❑ In Kepler it is possible to specify the size of shared memory banks
 - ❑ This changes the bank conflicts depending on your access pattern
- ❑ `cudaDeviceSharedMemConfig`
 - ❑ Options are;
 - ❑ `cudaSharedMemBankSizeDefault`: default size
 - ❑ `cudaSharedMemBankSizeFourByte`: 32 bit size banks
 - ❑ `cudaSharedMemBankSizeEightByte`: 64 bit size banks
 - ❑ Can improve performance of access for even strides.
- ❑ Not available in Maxwell +.