

Parallel Computing with GPUs

Performance


Part 1 – Memory Coalescing



The
University
Of
Sheffield.

Dr Paul Richmond


<http://paulrichmond.shef.ac.uk/teaching/COM4521/>



This Lecture (learning objectives)

❑ Memory Coalescing

- ❑ Explain the process of memory moving via cache lines
- ❑ Analyse the performance implications of strided access patterns
- ❑ Classify use cases of memory access with respect to performance




Coalesced Global Memory Access

- ❑ When memory is loaded/stored from global memory to L2 (and L1) it is moved in cache lines
 - ❑ If threads within a warp access global memory in irregular patterns this can cause increased movement (transactions) of data
- ❑ Coalesced access is where sequential threads in a warp access sequentially adjacent 4 byte words (e.g. `float` or `int` values).
 - ❑ Having coalesced access will reduce the number of cache lines moved and increase memory performance
 - ❑ This is one of the most important performance considerations of GPU memory usage!



Use of Memory Cache Levels




❑ L1 Cache


- ❑ 128B wide cache line transactions
- ❑ Normally used for thread local variables
- ❑ Can also be used for global loads (via read-only cache method from previous lecture)
- ❑ Some hardware has L1 cache for all memory reads
 - ❑ Always via L2 cache first
 - ❑ Compute **3.5**, 3.7 or 5.2 have opt in global L1 caching
 - ❑ Early Maxwell (Compute 5.0 cant opt in for L1 global loads)

❑ L2 Cache

- ❑ 32B wide cache line transactions
- ❑ **All** global and local memory pass through




L2 Coalesced Memory Access




```
__global__ void copy(float *odata, float* idata)
{
    int xid = blockIdx.x * blockDim.x + threadIdx.x;
    odata[xid] = idata[xid];
}
```

- ❑ Global memory always moves through L2
 - ❑ But not always through L1 depending on architecture
- ❑ In L2 cache line size is 32B
 - ❑ For a coalesced read/write within a warp, 4 transactions required
 - ❑ 100% memory bus speed




L2 Permuted Memory Access



☐ Permuted Access

- ☐ Within the cache line accesses can be permuted between threads
- ☐ No performance penalty

L2 Permuted Memory Access



☐ Permuted Access


☒ Permuted access within 128 byte segments is permitted

☐ Will NOT cause multiple loads

☐ Must not be permuted over the 128 byte boundary



L2 Offset Memory Access



```
__global__ void copy(float *odata, float* idata)
{
    int xid = blockIdx.x * blockDim.x + threadIdx.x + OFFSET;
    odata[xid] = idata[xid];
}
```



❑ If memory accesses are offset then parts of the cache line will be unused (shown in red) e.g.

❑ 5 transactions of 160B of which 128B is required: 80% utilisation

❑ Use thread block sizes of multiples of 32!


L2 Strided Memory Access

```
__global__ void copy(float *odata, float* idata)
{
    int xid = (blockIdx.x * blockDim.x + threadIdx.x) * STRIDE;
    odata[xid] = idata[xid];
}
```

❑ How many cache lines transactions for warp if $\text{STRIDE} = 2$?



L2 Strided Memory Access




```
__global__ void copy(float *odata, float* idata)
{
    int xid = (blockIdx.x * blockDim.x + threadIdx.x) * STRIDE;
    odata[xid] = idata[xid];
}
```



- ❑ Strided memory access can result in bad performance e.g.
 - ❑ A stride of 2 causes **8 transactions**: 50% useful memory bandwidth
 - ❑ As stride of >32 causes 32 transactions: ONLY 3.125% bus utilisation!
 - ❑ This is as bad as random access
 - ❑ Transpose data if it is stride-N

Degradation in Strided Access Performance




Array of Structures vs Structures of Arrays

❑ Array of Structures (AoS)

❑ Common method to store groups of data (e.g. points)

```
struct point {  
    float x, y, z;  
};  
__device__ struct point d_points[N];  
  
__global__ void manipulate_points()  
{  
    float x = d_points[blockIdx.x*blockDim.x + threadIdx.x].x;  
    float y = d_points[blockIdx.x*blockDim.x + threadIdx.x].y;  
    float z = d_points[blockIdx.x*blockDim.x + threadIdx.x].z;  
  
    func(x, y, z);  
}
```

Is this a good kernel?



Array of Structures vs Structures of Arrays

❑ Array of Structures (AoS)

❑ Common method to store groups of data (e.g. points)

```
struct point {  
    float x, y, z;  
};  
__device__ struct point d_points[N];  
  
__global__ void manipulate_points()  
{  
    float x = d_points[blockIdx.x*blockDim.x + threadIdx.x].x;  
    float y = d_points[blockIdx.x*blockDim.x + threadIdx.x].y;  
    float z = d_points[blockIdx.x*blockDim.x + threadIdx.x].z;  
  
    func(x, y, z);  
}
```



Is this a good kernel? **No: Stride of 3 has only 33% memory bandwidth**


Array of Structures vs Structures of Arrays

❑ An Alternative: Structure of Arrays (SoA)

```
struct points {  
    float x[N], y[N], z[N];  
};  
  
__device__ struct points d_points;  
  
__global__ void manipulate_points()  
{  
    float x = d_points.x[blockIdx.x*blockDim.x + threadIdx.x];  
    float y = d_points.y[blockIdx.x*blockDim.x + threadIdx.x];  
    float z = d_points.z[blockIdx.x*blockDim.x + threadIdx.x];  
  
    func(x, y, z);  
}
```



100% effective memory bandwidth



Summary

❑ Memory Coalescing

- ❑ Explain the process of memory moving via cache lines
- ❑ Analyse the performance implications of strided access patterns
- ❑ Classify use cases of memory access with respect to performance

❑ Next Lecture: The L1 Cache (and more on Coalescing)

