


# Parallel Computing with GPUs

## Warp Level CUDA and Atomics Part 3 – Atomics and Warp Operations



Dr Paul Richmond

<http://paulrichmond.shef.ac.uk/teaching/COM4521/>




# This Lecture (learning objectives)

## □ Atomics

- Present GPU atomic operations
- Demonstrate performance of GPU atomics and locks

## □ Warp Operations

- Present warp shuffle operations for communication of threads within a warp
- Give examples of warp operations such as sum and warp voting.





# What is wrong with the following

```
__global__ void max_kernel(int *a)
{
    __shared__ int max;

    int my_local = a[threadIdx.x + blockIdx.x*blockDim.x];

    if (my_local > max)
        max = my_local;
}
```



- ❑ More than one thread may try to modify max at the same time
  - ❑ Race condition

```
__global__ void max_kernel(int *a)
{
    __shared__ int max;

    int my_local = a[threadIdx.x + blockIdx.x*blockDim.x];

    if (my_local > max)
        max = my_local;
}
```





# Atomics

- ❑ Atomics are used to ensure correctness when concurrently reading and writing to a memory location (global or shared)


No need for assignment to a variable when using the atomic functions

```
__global__ void max_kernel(int *a)
{
    __shared__ int max;

    int my_local = a[threadIdx.x + blockIdx.x*blockDim.x];

    if (my_local > max)
        atomicMax(&max, my_local);
}
```

- ❑ No race condition
- ❑ Function supported in *most* hardware
  - ❑ Some older generation GPUs lack shared memory and floating point atomic etc.



# Atomic Functions and Locks



## ❑ An atomic function

- ❑ Must guarantee that an operation can complete without interference from any other thread
- ❑ Does not provide any guarantee of ordering or provide any synchronisation

## ❑ How can we implement critical sections?

```
__device__ int lock = 0;

__global__ void kernel() {
    bool need_lock = true;
    // get lock
    while (need_lock) {
        if (atomicCAS(&lock, 0, 1)==0) {
            //critical code section
            atomicExch(&lock, 0);
            need_lock = false;
        }
    }
}
```

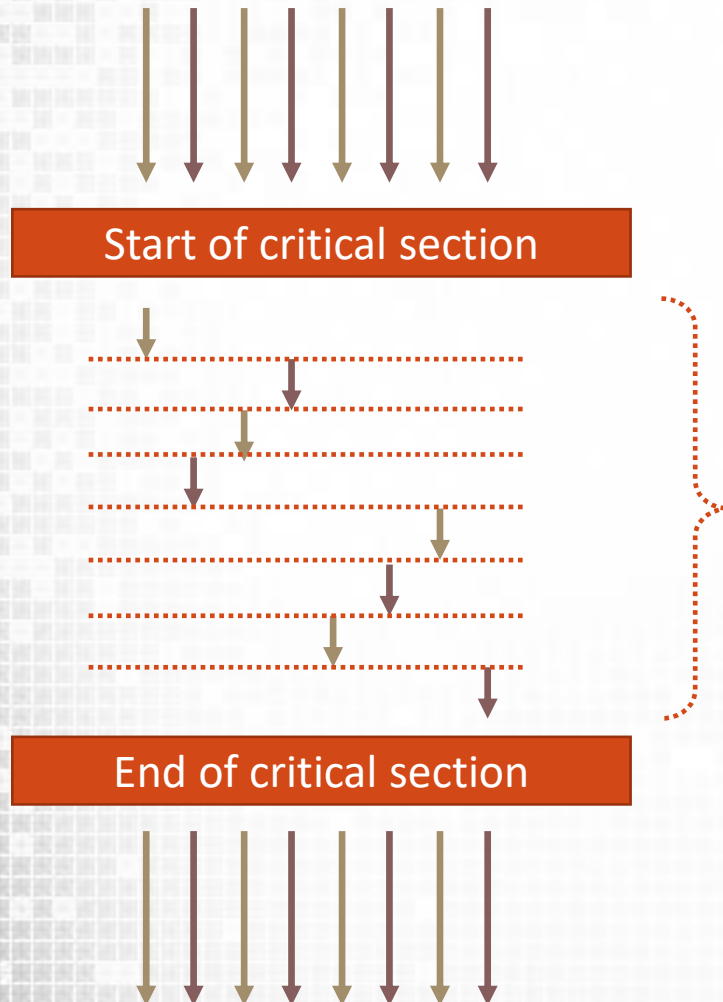
```
int atomicCAS(int* address, int compare, int val)
```

Performs the following in a single atomic transaction (atomic instruction)

```
*address = (*address == compare) ? val : *address;
```

Returning the old value at the address

# Serialisation




❑ What happens to performance when using atomics?

❑ In the case of the critical section example

❑ This is serialised for each thread accessing the shared value

❑ For the atomic CAS instruction access to the shared lock variable is serialised

❑ This is true of any atomic function or instruction in CUDA





# CUDA Atomic Functions / Instructions

❑ In addition to `atomicCAS` the following atomic functions/instructions are available

❑ Addition/subtraction

❑ E.g. `int atomicAdd(int* address, int val)` – add `val` to integer at `address`

❑ Exchange

❑ Exchange a value with a new value

❑ Increment/Decrement

❑ Minimum and Maximum

❑ Variants of atomic functions

❑ 64 bit integer and double versions available in Pascal (Compute 6.0)


❑ See docs: <https://docs.nvidia.com/cuda/cuda-c-programming-guide/index.html#atomic-functions>






# Local vs Global Atomics

Kepler



Maxwell



## ❑ Image histogram example

❑ Accumulation of colour values for images


❑ Entropy: measure of the level of disorder (lower entropy == higher contention)

❑ <https://devblogs.nvidia.com/parallelforall/gpu-pro-tip-fast-histograms-using-shared-atomics-maxwell/>




# Warp Shuffle

- ❑ For moving/comparing data between threads in a block it is possible to use Shared Memory (SM)
- ❑ For moving/comparing data between threads in a warp (known as lanes in this context) it is possible to use a *warp shuffle* (SHFL)
  - ❑ Direct exchange of information between two threads
    - ❑ Can replace atomics
    - ❑ Should never depend on conditional execution!
  - ❑ Does not require SM
    - ❑ Always faster than SM equivalent
  - ❑ Implicit synchronisation (no need for `__syncthreads`)
    - ❑ EXCEPT on Volta+ hardware (use `__syncwarp`)
  - ❑ Works by allowing threads to read another threads registers
  - ❑ <https://docs.nvidia.com/cuda/cuda-c-programming-guide/index.html#warp-shuffle-functions>




# Shuffle Variants

`__shfl_sync()`




Shuffled between  
any two index  
threads

`__shfl_up_sync()`




Shuffles to  $n^{\text{th}}$  right  
neighbour wrapping  
indices (in this case  
 $n=2$ )

`__shfl_down_sync()`




Shuffles to  $n^{\text{th}}$  left  
neighbour wrapping  
indices (in this case  
 $n=2$ )

`__shfl_xor_sync()`




Butterfly (XOR)  
exchange shuffle  
pattern





# Shuffle function arguments

- ❑ `T __shfl_sync(unsigned mask, T var, int srcLane, int width=warpSize);`
- ❑ `T __shfl_up_sync(unsigned mask, T var, unsigned int delta, int width=warpSize);`
  - ❑ delta is the n step used for shuffling
- ❑ `T __shfl_down_sync(unsigned mask, T var, unsigned int delta, int width=warpSize);`
  - ❑ Source lane determined by bitwise XOR with laneMask
- ❑ Mask is a bit mask for the warp to indicate which threads participate
- ❑ T can be `int`, `unsigned int`, `long`, `unsigned long`, `long long`, `unsigned long long`, `float` or `double`
- ❑ Optional width argument
  - ❑ Must be a power of 2 and less than or equal to warp size
  - ❑ If smaller than warp size each subsection acts independently (own wrapping)






# Shuffle Warp Sum Example (down)


```
__global__ void sum_warp_kernel_shfl_down(int *a)
{
    int local_sum = a[threadIdx.x + blockIdx.x*blockDim.x];

    for (int offset = WARP_SIZE / 2; offset>0; offset /= 2)
        local_sum += __shfl_down(local_sum, offset);

    if (threadIdx.x%32 == 0)
        printf("Warp max is %d", local_sum)
}
```



Warp sum in `threadIdx.x%32==0`




# Shuffle Warp Sum Example (xor)


```
__global__ void sum_warp_kernel_shfl_xor(int *a)
{
    int local_sum = a[threadIdx.x + blockIdx.x*blockDim.x];

    for (int mask = WARP_SIZE / 2; mask>0; mask /= 2)
        local_sum += __shfl_xor(local_sum, mask);

    if (threadIdx.x%32 == 0)
        printf("Warp max is %d", local_sum)
}
```



Warp sum in all threads



# Warp Voting

- ❑ Warp shuffles allow data to be exchanged between threads in a warp
- ❑ Warp voting allows threads to test a condition across all threads in a warp
  - ❑ `int all(condition)`
    - ❑ True if the condition is met by all threads in the warp
  - ❑ `int any(condition)`
    - ❑ True if any thread in warp meets condition
  - ❑ `unsigned int ballot(condition)`
    - ❑ Sets the  $n^{\text{th}}$  bit of the return value based on the  $n^{\text{th}}$  threads condition value
- ❑ All warp voting functions are single instruction and act as barrier
  - ❑ Only active threads participate, does not block like `syncthreads`




# Warp Voting Example

```
__global__ void voteAllKernel(unsigned int *input, unsigned int *result)
{
    int i = threadIdx.x + blockIdx.x*blockDim.x;
    int j = i % WARP_SIZE;

    int vote_result = all(input[i]);

    if (j==0)
        result[i / WARP_SIZE] = vote_result;
```

- ❑ For each first thread in the warp calculate if all threads in the warp have `true` valued input
- ❑ Save the warp vote to a compact array
  - ❑ A reduction of factor 32






# Global Communication Summary

- ❑ Shared memory is per thread block
- ❑ Shuffles and voting for warp level
- ❑ Atomics can be used for some global (grid wide) operations
- ❑ What about general global communication?
  - ❑ Not possible within a kernel (*except in Volta – not covered*)!
  - ❑ Remember a grid may not be entirely in flight on the device
  - ❑ Can be enforced by finishing the kernel

```
step1 <<<grid, blk >>>(input, step1_output);  
// step1_output can safely be used as input for step2  
step2 <<<grid, blk >>>(step1_output, step2_output);
```




# Summary

## ❑ Atomics

- ❑ Present GPU atomic operations
- ❑ Demonstrate performance of GPU atomics and locks

## ❑ Warp Operations

- ❑ Present warp shuffle operations for communication of threads within a warp
- ❑ Give examples of warp operations such as sum and warp voting.



# Acknowledgements and Further Reading

- ❑ Predication: <http://docs.nvidia.com/cuda/parallel-thread-execution/index.html#predicated-execution>
- ❑ Shuffling: <http://on-demand.gputechconf.com/gtc/2013/presentations/S3174-Kepler-Shuffle-Tips-Tricks.pdf>
- ❑ Volta: <https://devblogs.nvidia.com/cuda-9-features-revealed/>

