

Parallel Computing with GPUs

Parallel Patterns Part 3 – Scan



Dr Paul Richmond

<http://paulrichmond.shef.ac.uk/teaching/COM4521/>



This Lecture (learning objectives)


□ Scan

- Give motivating examples of parallel prefix sum (scan)
- Describe the serial and parallel approaches towards scan
- Compare block level and atomic approaches to the parallel prefix sum algorithm




What is scan?


□ Consider the following ...



Output variable numbers of values per thread



Remove empty elements from array (compact)




Split elements from array based on condition (split)




What is scan?


❑ Consider the following ...



Output variable numbers of values per thread



Remove empty elements from array (compact)



Split elements from array based on condition (split)

❑ Each has the same problem

❑ Not even considered for sequential programs!


❑ Where to write output in parallel?



Parallel Prefix Sum (scan)

❑ Where to write output in parallel?

❑ Each threads needs to know the output location(s) it can write to avoid conflicts.



Thread/Read index

Output/Write index – running sum of binary output state

Sparse data

Compacted data

❑ The solution is a parallel prefix sum (or scan)

❑ Given the inputs $A = [a_0, a_1, \dots, a_{n-1}]$ and binary associate operator \oplus

❑ $\text{Scan}(A) = [0, a_0, (a_0 \oplus a_1), \dots, (a_0 \oplus a_1 \oplus \dots \oplus a_{n-1})]$



Serial Parallel Prefix Sum Example

❑ E.g. Given the input and the addition operator

❑ $A = [2, 6, 2, 4, 7, 2, 1, 5]$

❑ $\text{Scan}(A) = [0, 2, 2+6, 2+6+2, 2+6+2+4, \dots]$

❑ $\text{Scan}(A) = [0, 2, 8, 10, 14, 21, 23, 24]$

❑ More generally a serial implementation of an additive scan using a running sum looks like...

```
int A[8] = { 2, 6, 2, 4, 7, 2, 1, 5 };
int scan_A[8];
int running_sum = 0;
for (int i = 0; i < 8; ++i)
{
    scan_A[i] = running_sum;
    running_sum += A[i];
}
```



Serial Scan for Compaction



```
int Input[8] = { 2, 0, 5, 6, 3, 0, 1, 0 };
int A[8] =      { 2, 0, 5, 6, 3, 0, 1, 0 };
int scan_A[8];
int output[5]
int running_sum = 0;

for (int i = 0; i < 8; ++i){
    A[i] = Input>0;
}

for (int i = 0; i < 8; ++i){
    scan_A[i] = running_sum;
    running_sum += A[i];
}

for (int i = 0; i < 8; ++i){
    int input = Input[i];
    if (input > 0){
        int idx = scan[i];
        output[idx] = input;
    }
}
```

```
// generate scan input
// A = {1, 0, 1, 1, 1, 0, 1, 0}
```

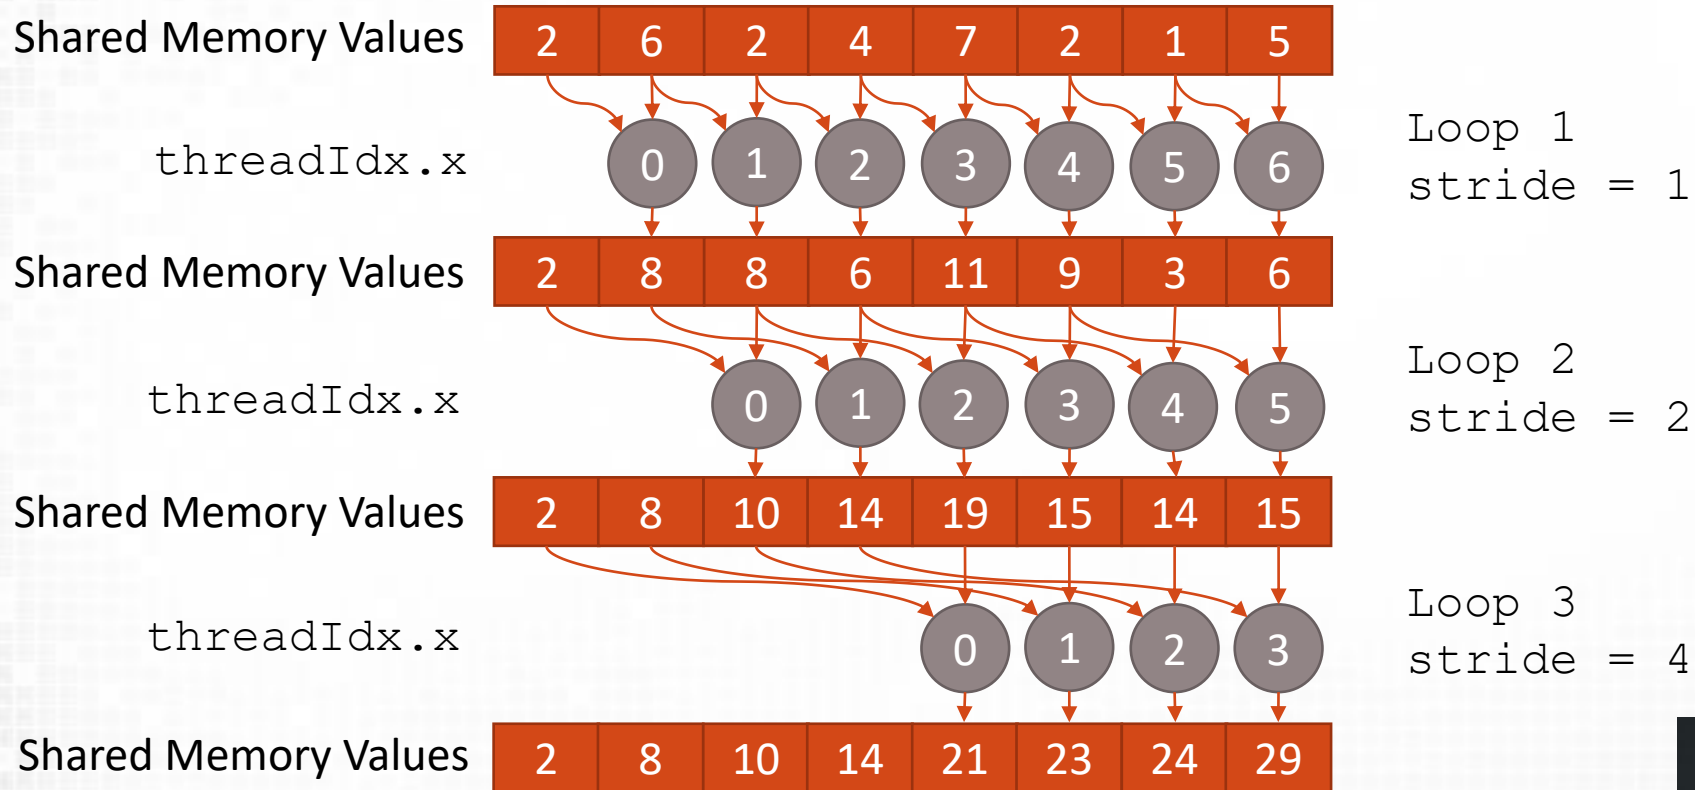
```
// scan
// scan_A = {0, 1, 1, 2, 3, 4, 4, 5}
```

We could test either Input[i] or A[i] to find empty values

```
// scattered write
// output = {2, 5, 6, 3, 1}
```


Parallel Local (Shared Memory) Scan

After $\log(N)$ loops each sum has local plus preceding $2^n - 1$ values

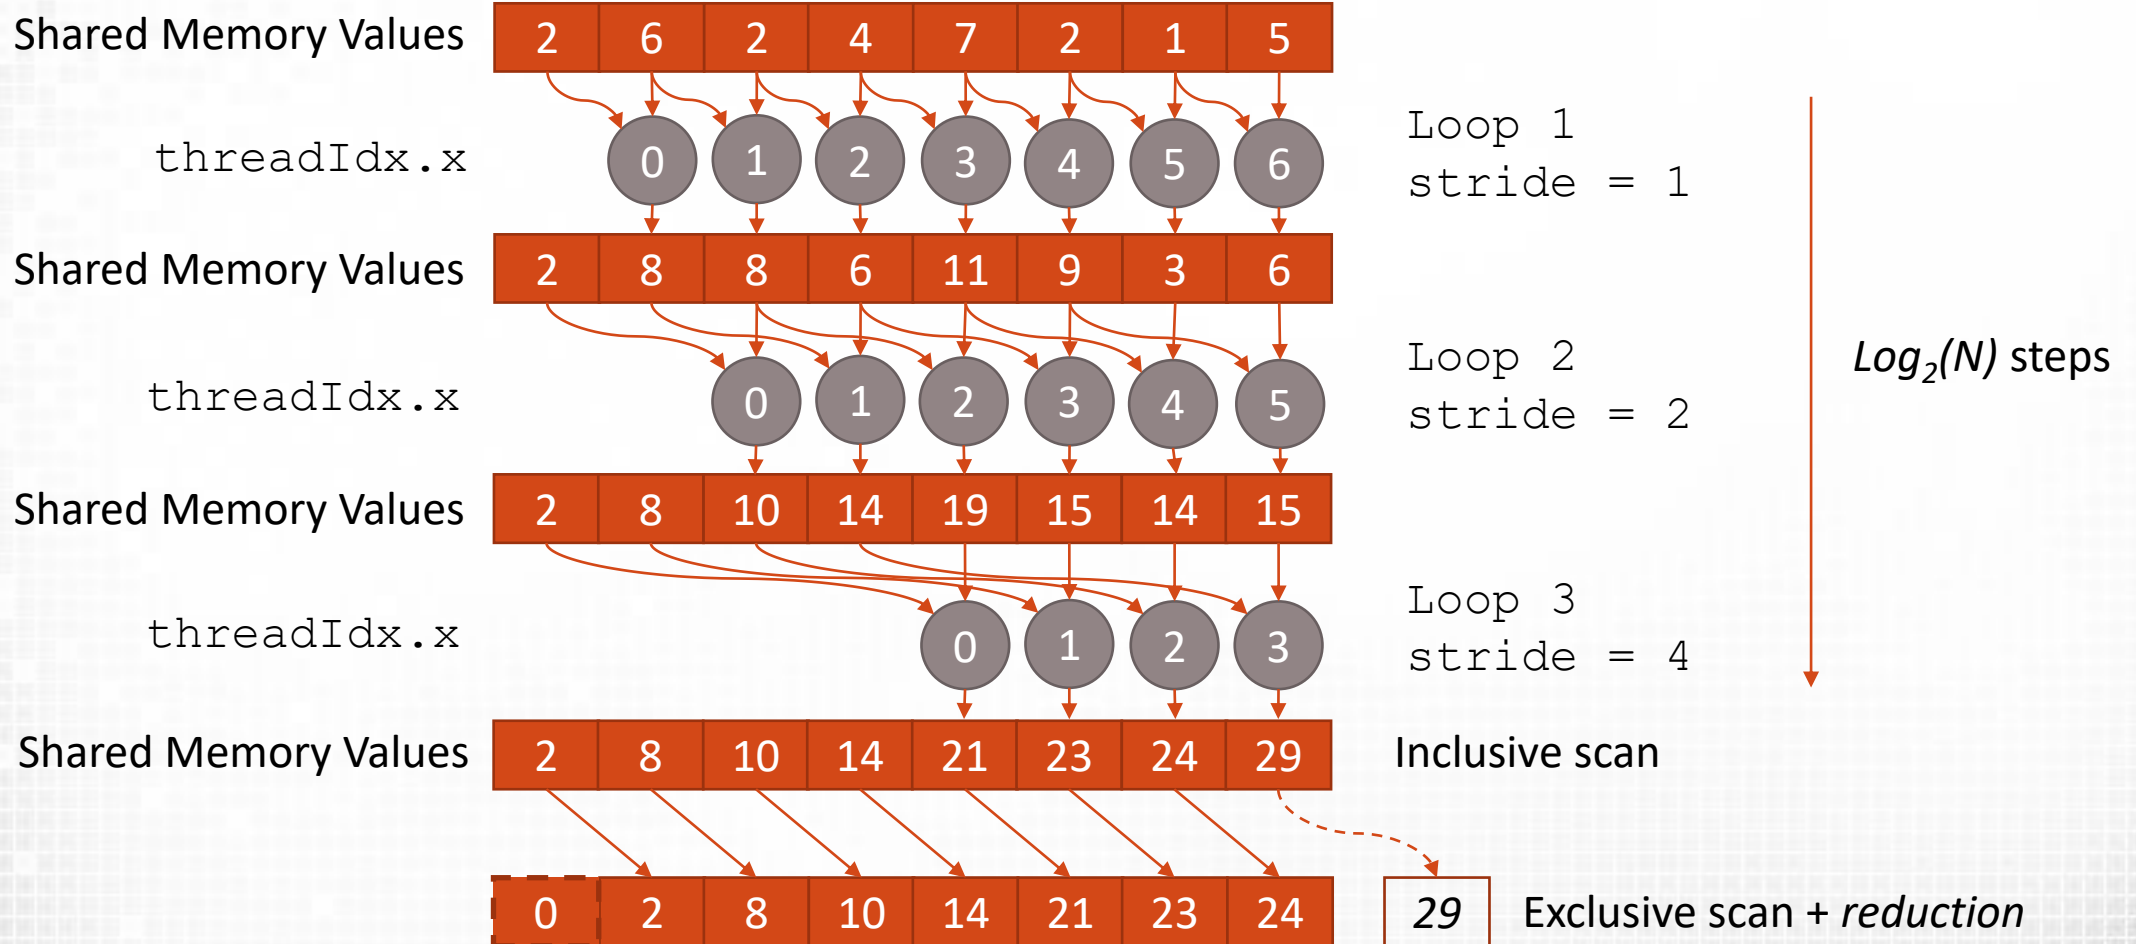


$\log_2(N)$ steps

Inclusive Scan



Parallel Local Scan



Implementing Local Scan with Shared Memory

```
__global__ void scan(float *input) {  
    extern __shared__ float s_data[];  
    s_data[threadIdx.x] = input[threadIdx.x + blockDim.x];  
  
    for (int stride = 1; stride < blockDim.x; stride <= 1) {  
        __syncthreads();  
        float s_value = (threadIdx.x >= stride) ? s_data[threadIdx.x - stride] : 0;  
        __syncthreads();  
        s_data[threadIdx.x] += s_value;  
    }  
  
    //something with global results?  
}
```


- ❑ No bank conflicts (stride of 1 between threads)
- ❑ Synchronisation required between read and write



Implementing Local Scan (at warp level)

```
__global__ void scan(float *input) {  
    __shared__ float s_data[32];  
    float val1, val2;  
  
    val1 = input[threadIdx.x + blockIdx.x*blockDim.x];  
  
    for (int s = 1; s < 32; s <= 1) {  
        val2 = __shfl_up(val1, s);  
        if (threadIdx.x % 32 >= s)  
            val1 += val2;  
    }  
  
    //store warp level results}
```

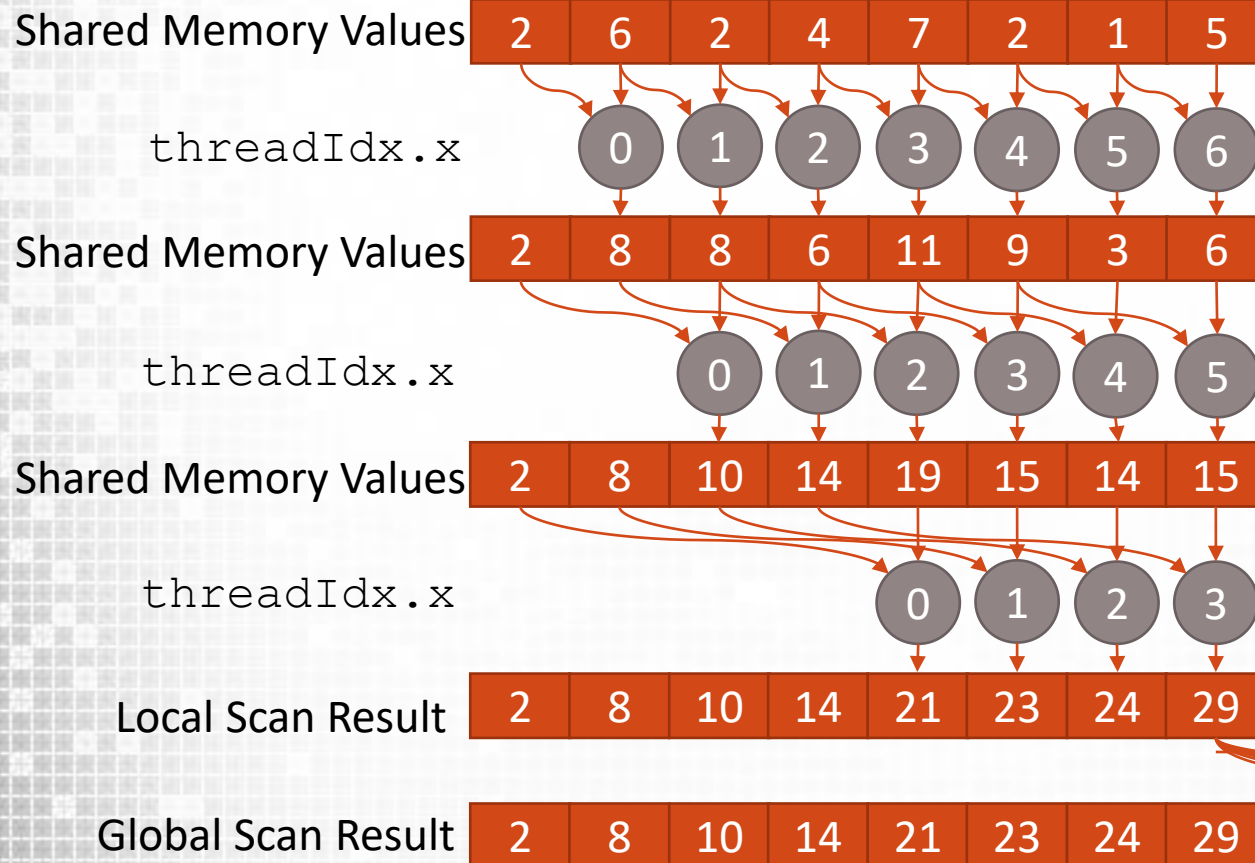
- ❑ Exactly the same as the block level technique but at warp level
- ❑ Warp prefix sum is in `threadIdx.x % 32 == 31`
- ❑ Either use shared memory to reduce between warps
 - ❑ Or consider the following global scan approaches.



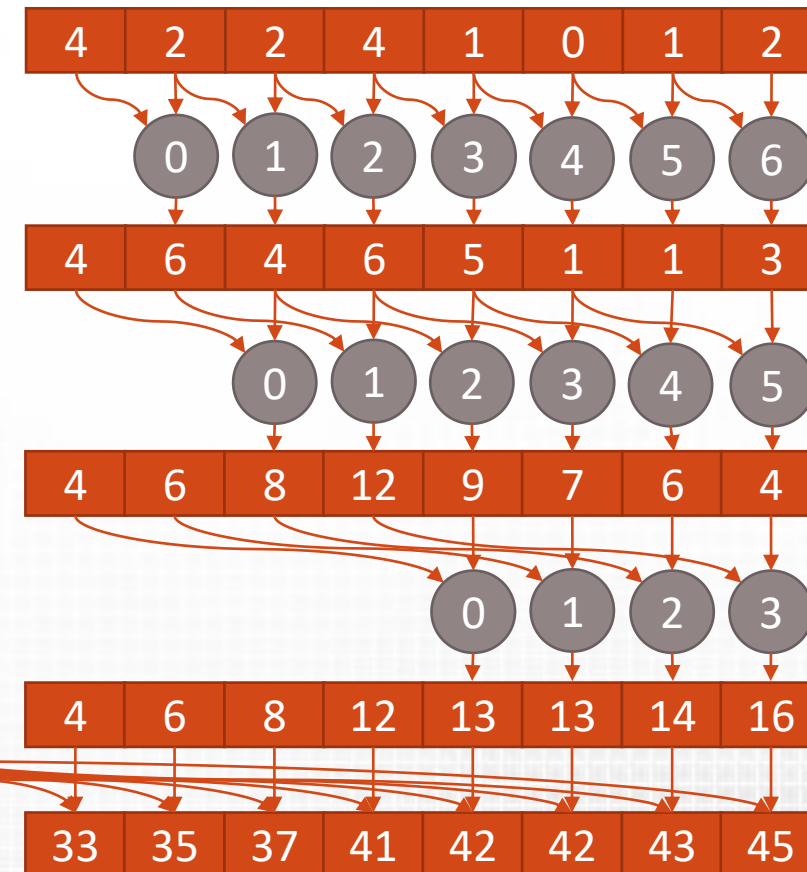
Implementing scan at Grid Level



Thread Block 1




Thread Block 2



Implementing scan at Grid Level

- ❑ Same problem as reduction when scaling to grid level
 - ❑ Each block is required to add the reduction value from proceeding blocks

- ❑ Global scan therefore requires either;
 1. Recursive scan kernel on results of local scan
 - ❑ Additional kernel to add sums of proceeding blocks
 2. **Atomic Increments (next slides)**
 - ❑ Increment a counter for block level results
 - ❑ Additional kernel to add sums of proceeding blocks to each value



Global Level Scan (Atomics Part 1)



```
__device__ block_sums[BLOCK_DIM];

__global__ void scan(float *input, float *local_result) {
    extern __shared__ float s_data[];
    s_data[threadIdx.x] = input[threadIdx.x + blockIdx.x*blockDim.x];

    for (int stride = 1; stride<blockDim.x; stride<=<1) {
        __syncthreads();
        float s_value = (threadIdx.x >= stride) ? s_data[threadIdx.x - stride] : 0;
        __syncthreads();
        s_data[threadIdx.x] += s_value;
    }

    //store local scan result to each thread
    local_result[threadIdx.x + blockIdx.x*blockDim.x] = s_data[threadIdx.x];

    //atomic store to all proceeding block totals (e.g. blocks after this block)
    if (threadIdx.x == 0){
        for (int i=blockIdx.x+1; i<gridDim.x; i++)
            atomicAdd(&block_sums[i], s_data[blockDim.x-1]);
    }
}
```

Global Level Scan (Atomics Part 2)

- ❑ After completion of the first kernel, block sums are all synchronised
- ❑ Use first thread in block to load block total into shared memory
- ❑ Increment local result


```
__device__ block_sums[BLOCK_DIM];

__global__ void scan_update(float *local_result, float *global_result) {
    extern __shared__ float block_total;
    int idx = threadIdx.x + blockIdx.x*blockDim.x;

    if (threadIdx.x == 0)
        block_total = block_sums[blockIdx.x];

    __syncthreads();


    global_result[idx] = local_result[idx]+block_total;
}
```



Summary

□ Scan

- Give motivating examples of parallel prefix sum (scan)
- Describe the serial and parallel approaches towards scan
- Compare block level and atomic approaches to the parallel prefix sum algorithm



Acknowledgements and Further Reading

- ❑ <https://devblogs.nvidia.com/paralleforall/faster-parallel-reductions-kepler/>

- ❑ All about application of warp shuffles to reduction

- ❑ https://stanford-cs193g-sp2010.googlecode.com/svn/trunk/lectures/lecture_6/parallel_patterns_1.ppt

- ❑ Scan material based loosely on this lecture

- ❑ http://docs.nvidia.com/cuda/samples/6_Advanced/reduction/doc/reduction.pdf

- ❑ Reduction material is based on this fantastic lecture by Mark Harris (NVIDIA)

