

Parallel Computing with GPUs

Sorting and Libraries Part 1 – Sorting



Dr Paul Richmond

<http://paulrichmond.shef.ac.uk/teaching/COM4521/>




This Lecture (learning objectives)

❑ Sorting Networks

- ❑ Demonstrate the use of a sorting network to achieve parallel sorting
- ❑ Compare sorting networks with serial sorting approaches

❑ Merge and Bitonic Sort

- ❑ Present the merge sort and considers its performance implications
- ❑ Identify performance features of bitonic sorting



Serial Sorting Examples

❑ Insertion Sort

❑ Insert a new element into a sorted list.

❑ E.g. [1 6 3 4 2 5]

❑ [1] -> [1 6] -> [1 3 6] -> [1 3 4 6] -> [1 2 3 4 6] -> [1 2 3 4 5 6]

❑ Bubble Sort

❑ Exchange and Sweep to compare each pair of adjacent elements

❑ $O(n^2)$ worst-case and average case, $O(n)$ best case.

❑ E.g. [1 6 3 4 2 5]

❑ [1 6 3 4 2 5] -> [1 3 6 4 2 5] -> [1 3 4 6 2 5] -> [1 3 4 2 6 5] -> [1 3 4 2 5 6]

❑ [1 3 2 4 5 6]

❑ [1 2 3 4 5 6]



Serial Sorting Examples

❑ Insertion Sort

❑ Insert a new element into a sorted list.

❑ E.g. [1 6 3 4 2 5]

❑ [1] -> [1 6] -> [1 3 6] -> [1 3 4 6] -> [1 2 3 4 6] -> [1 2 3 4 5 6]

❑ Bubble Sort

❑ Exchange and Sweep to compare each pair of adjacent elements

❑ $O(n^2)$ worst-case and average case, $O(n)$ best case.

❑ E.g. [1 6 3 4 2 5]

❑ [1 6 3 4 2 5] -> [1 3 6 4 2 5] -> [1 3 4 6 2 5] -> [1 3 4 2 6 5] -> [1 3 4 2 5 6]

❑ [1 3 2 4 5 6]

❑ [1 2 3 4 5 6]



Classifying Sort Techniques/Implementations

☐ Data driven

- ☐ Each step of the algorithm depends on the previous step version
- ☐ Highly serial


☐ Data independent

- ☐ The algorithm performs fixed steps and does not change its processing based on data
- ☐ Well suited to parallel implementations
- ☐ Can be expressed as a sorting network...



Sorting Networks

- ❑ A sorting network is a comparator network that sorts all input sequences
- ❑ Following the same execution of stages
- ❑ Consider the previous Bubble Sort [1 6 3 4 2 5]




[1 6 3 4 2 5] -> [1 3 6 4 2 5] -> [1 3 4 6 2 5] -> [1 3 4 2 6 5] -> [1 3 4 2 5 6]
[1 3 4 2 5 6] -> [1 3 4 2 5 6] -> [1 3 2 4 5 6] -> [1 3 2 4 5 6]
[1 3 2 4 5 6] -> [1 2 3 4 5 6] -> [1 2 3 4 5 6]
[1 3 2 4 5 6] -> [1 2 3 4 5 6]
[1 2 3 4 5 6]

Sweeps

Not considered

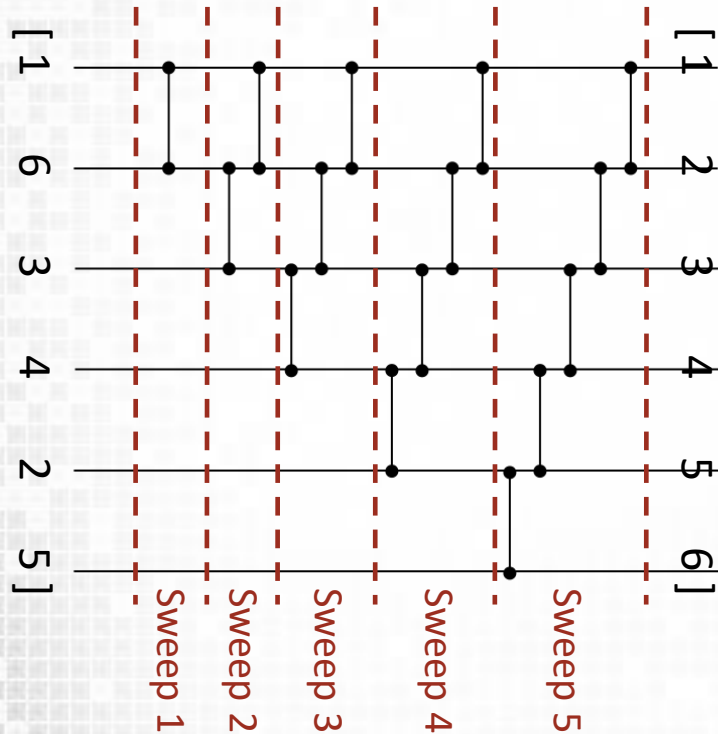
Compared not swapped

Compared and swapped



Sorting Networks

□ And Insertion Sort...




[1 6 3 4 2 5]
[1 3 6 4 2 5] → [1 3 6 4 2 5]
[1 3 4 6 2 5] → [1 3 4 6 2 5] → [1 3 4 6 2 5]
[1 3 4 2 6 5] → [1 3 2 4 6 5] → [1 2 3 4 6 5] → [1 2 3 4 6 5]
[1 2 3 4 5 6] → [1 2 3 4 5 6] → [1 2 3 4 5 6] → [1 2 3 4 5 6] → [1 2 3 4 5 6]

Sweeps


Not considered

Compared not swapped


Compared and swapped




Parallel Sorting Networks



Bubble



Insertion




❑ Parallel Bubble and Insertion sorting network is still not very efficient

❑ $2n - 3$ sweeps

❑ $n(n - 1)/2$ comparisons - $O(n^2)$ complexity



[1 6 3 4 2 5]
[1 3 6 4 2 5]
[1 3 4 6 2 5]
[1 3 4 2 6 5]
[1 3 2 4 5 6]
[1 2 3 4 5 6]
[1 2 3 4 5 6]
[1 2 3 4 5 6]

Sweeps = 9






Merge Sort

- ❑ To reduce the $O(n^2)$ overhead we need a better sorting network
- ❑ The odd-even merge sort network (for power of 2 n)
 - ❑ Sort all odd and even keys separately and then merge m values of a stage
 - ❑ Merge a sorted sequence of elements on lines $\langle a_1, \dots, a_n \rangle$ with those on lines $\langle a_{n+1}, \dots, a_{2n} \rangle$
 - ❑ Each merge requires $\log(n)$ passes
 - ❑ Total complexity of $O(n \log(n^2) + \log(n))$






Merge Sort


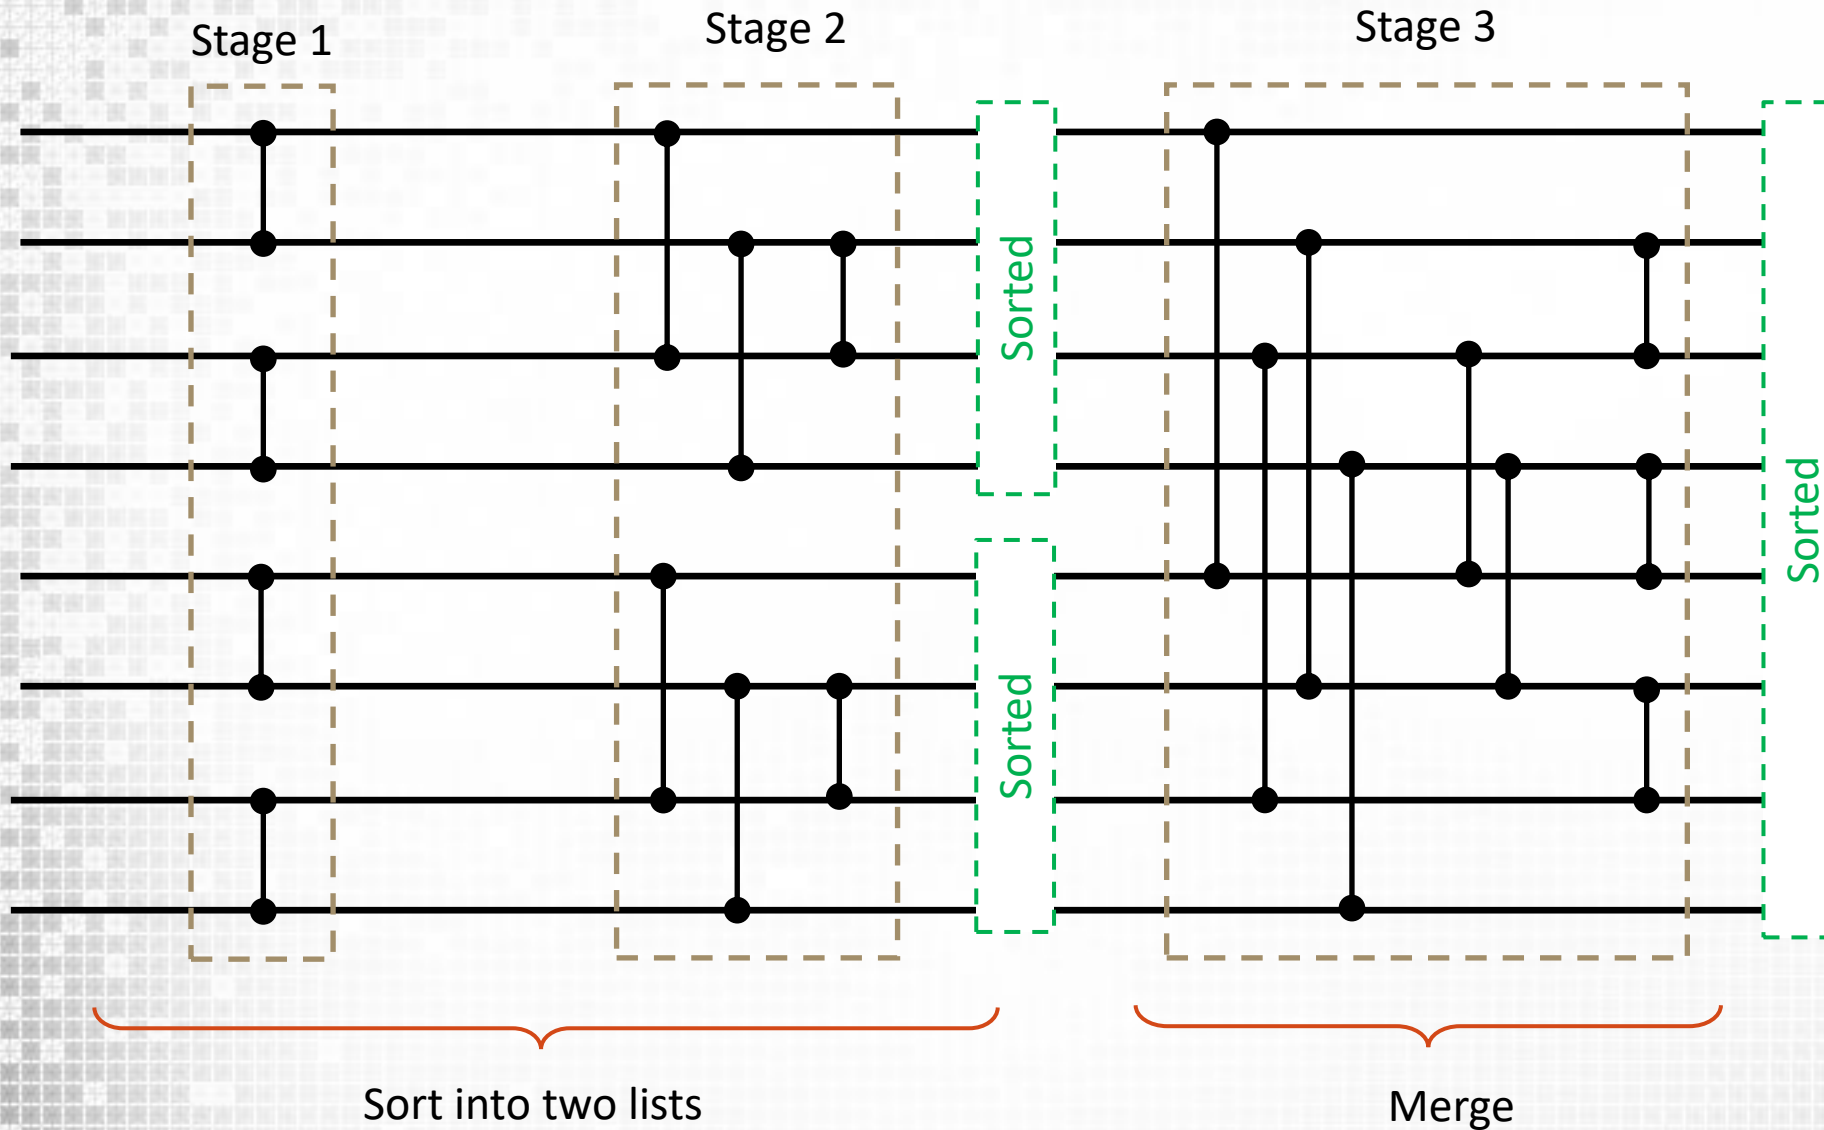
- ❑ To reduce the $O(n^2)$ overhead we need a better sorting network
- ❑ The odd-even merge sort network (for power of 2 n)
 - ❑ Sort all odd and even keys separately and then merge m values of a stage
 - ❑ Merge a sorted sequence of elements on lines $\langle a_1, \dots, a_n \rangle$ with those on lines $\langle a_{n+1}, \dots, a_{2n} \rangle$
 - ❑ Each merge requires $\log(n)$ passes
 - ❑ Total complexity of $O(n \log(n^2) + \log(n))$




Merging of two sorted sequences ($n=4$)



Full Merge Sorting (n=4)





Full Merge Sorting (n=4) example



Input	Stage 1		Stage 2		Stage 3						Output
8	1		1		1						1
1	8			5 3			3			2	2
5	3		3		5			2		3	3
3	5			8				8	4	4	4
6	2		2		2			5		5	5
2	6			6 4			4		8	6	6
4	4		4		6					8	8
9	9			9				9			9

Limitations of Merge Sort

- ❑ What is potentially wrong with a merge sort GPU implementation?
 - ❑ Irregular memory accesses
 - ❑ Not all values are compared in each pass (uneven workload per thread)





Solution: Bitonic Sort

□ Bitonic sorting network

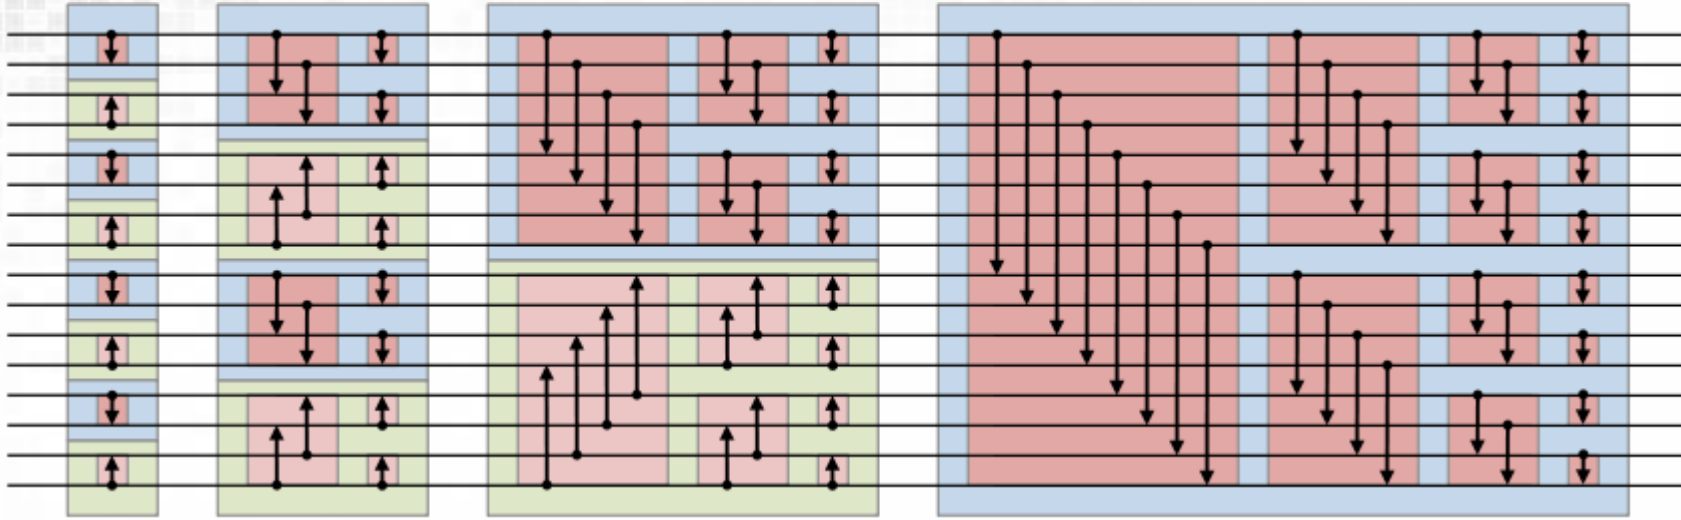
- Iterative splitting and merging of inputs into increasing large bitonic sequences

- A sequence is bitonic if


- There is an i , such that $a_0 \dots, a_i$ is monotonically increasing and $a_i \dots, a_n$ is monotonically decreasing



Bitonic Sorting Network



- ❑ Sorting and Merging increasing large bitonic sequences
 - ❑ When $n = 2^k$ there are k levels with $\frac{n}{2}$ comparisons each
- ❑ GPU Implementation
 - ❑ Regular access strides :-)
 - ❑ Efficiently balanced workload :-)
 - ❑ Requires multiple kernel launches to merge over $n > \text{block size}$



Summary

❑ Sorting Networks

- ❑ Demonstrate the use of a sorting network to achieve parallel sorting
- ❑ Compare sorting networks with serial sorting approaches

❑ Merge and Bitonic Sort

- ❑ Present the merge sort and considers its performance implications
- ❑ Identify performance features of bitonic sorting

❑ Next: Libraries and Thrust

