

Parallel Computing with GPUs

Sorting and Libraries Part 2 – Libraries



Dr Paul Richmond

<http://paulrichmond.shef.ac.uk/teaching/COM4521/>



This Lecture (learning objectives)


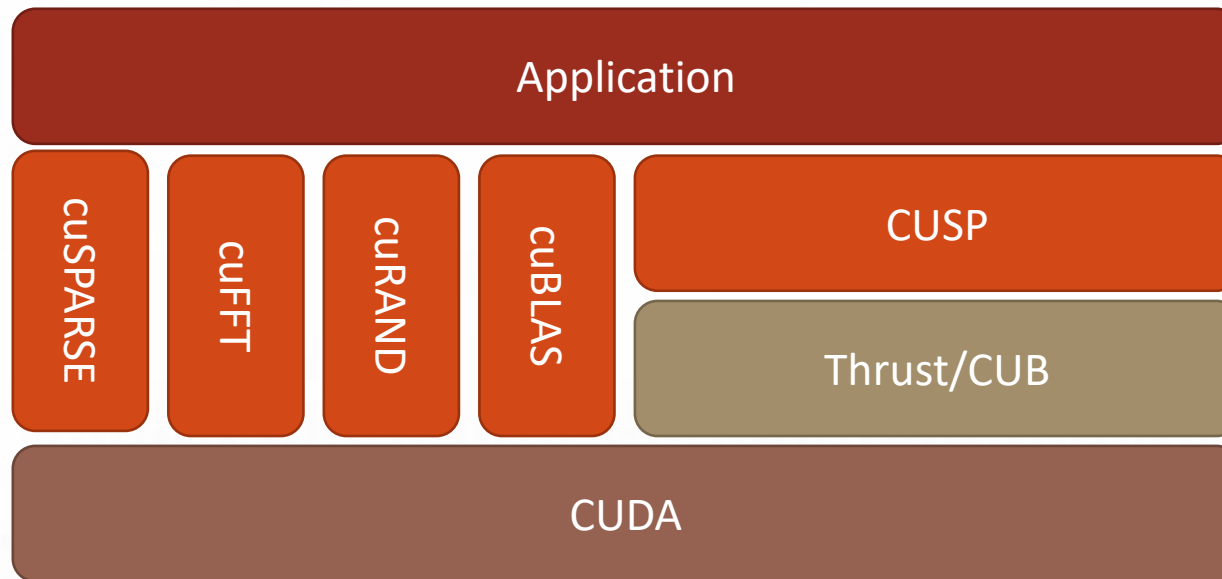
❑ Libraries and Thrust

- ❑ Describe the purpose of CUDA libraries
- ❑ Demonstrate Thrust containers for data storage
- ❑ Explain the relationship between raw pointers and Thrust iterators
- ❑ Give example of Thrust algorithms



CUDA libraries

- ❑ Abstract CUDA model away from programmer
- ❑ Highly optimised implementations of common tools
 - ❑ Mainly focused on linear algebra




Thrust

❑ Template Library for CUDA

- ❑ Implements many parallel primitives (scan, sort, reduction etc.)
- ❑ Part of standard CUDA release
- ❑ Level of Abstraction which hides kernels, mallocs and memcpy's

❑ Designed for C++ programmers

- ❑ Similar in design and operation as the C++ Standard Template Library (STL)
- ❑ Only a small amount of C++ required..



Thrust containers

❑ Thrust uses only high level *vector* containers

❑ `host_vector`: on host

❑ `device_vector`: on GPU

❑ Other STL containers include

❑ `queue`

❑ `list`

❑ `tack`

❑ `queue`

❑ `priority_queue`

❑ `set`

❑ `multiset`

❑ `map`

❑ `multimap`

❑ `bitset`

❑ STL containers can be used to initialise a Thrust vector


```
#include <thrust/host_vector.h>
#include <thrust/device_vector.h>

int main()
{
    //create a vector on the host
    thrust::host_vector<int> h_vec(10);

    //create a vector on the device
    thrust::device_vector<int> d_vec = h_vec;

    //device data manipulated directly from host
    for (int i = 0; i < 10; i++)
        d_vec[i] = i;

    //vector memory automatically released
    return 0;
}
```




Thrust Iterators

- ❑ They point to regions of a vector
- ❑ Can be used like pointers
 - ❑ Explicit cast when dereferencing very important

```
thrust::device_vector<int>::iterator begin = d_vec.begin();  
thrust::device_vector<int>::iterator end = d_vec.end();  
printf("d_vec at begin=%d", (int)*begin);  
begin++; //move on a single position  
printf("d_vec at begin++=%d", (int)*begin);  
*end = 88;  
printf("d_vec at end=%d", (int)*end);
```

```
d_vec at begin=0  
d_vec at begin++=1  
d_vec at end=88
```



Thrust Iterators



❑ Can be converted to a raw pointer

```
int * d_ptr = thrust::raw_pointer_cast(begin);  
int * d_ptr = thrust::raw_pointer_cast(begin[0]);  
  
kernel<BLOCKS, TPB>(d_ptr);
```

❑ Raw pointers can be used in Thrust

❑ BUT not exactly the same as a vector

```
int* d_ptr;  
cudaMalloc((void**) &d_ptr, N);  
  
thrust::device_ptr<int> d_vec = thrust::device_pointer_cast(d_ptr);  
//or  
thrust::device_ptr<int> d_vec = thrust::device_ptr<int>(d_ptr)
```

Thrust Algorithms

❑ Transformations

- ❑ Application of a function to each element within the range of a vector

❑ Reduction

- ❑ Reduction of a set of values to a single value using binary associative operator
- ❑ Can also be used to count occurrences of a value

❑ Prefix Sum

- ❑ Both inclusive and exclusive scans

❑ Sort

- ❑ Can sort keys or key value pairs

❑ Binary Search

- ❑ Position of a target value



Thrust Transformations

□ Some examples of the many transformations


```
//copy a vector (or part of a vector) to another vector
thrust::copy(d_vec.begin(), d_vec.begin() + 10, d_vec_cpy.begin());

//fill a vector with a value
thrust::fill(d_vec.begin(), d_vec.begin() + 10, 0);

//rand is a predefined Thrust generator
thrust::generate(d_vec.begin(), d_vec.begin() + 10, rand);

// fill d_vec with {0, 1, 2, 3, 4, 5, 6, 7, 8, 9}
thrust::sequence(d_vec.begin(), d_vec.begin() + 10);

//all occurrences of the value 1 are replaced with the value 10
thrust::replace(d_vec.begin(), d_vec.end(), 1, 10);
```



Thrust Algorithms

□ Either in-place or to output vector

```
thrust::device_vector<int> d_vec(10);
thrust::device_vector<int> d_vec_out(10);

//fill d_vec with {0, 1, 2, 3, 4, 5, 6, 7, 8, 9}
thrust::sequence(d_vec.begin(), d_vec.begin() + 10);


//inclusive scan to output vector
thrust::inclusive_scan(d_vec.begin(), d_vec.end(), d_vec_out.begin());

//inclusive scan in place
thrust::inclusive_scan(d_vec.begin(), d_vec.end(), d_vec.begin());

//generate random data (actually a transformation)
thrust::generate(d_vec.begin(), d_vec.end(), rand);

//sort in place
thrust::sort(d_vec.begin(), d_vec.end());

//sort data from a raw pointer (N is number of elements)
thrust::device_ptr<int> dt_ptr = thrust::device_pointer_cast(d_a_ptr);
thrust::sort(dt_ptr, dt_ptr+N);
```



Custom Transformations

```
thrust::device_vector<int> d_vec(10);
thrust::device_vector<int> d_vec_out(10);


//fill d_vec with {0, 1, 2, 3, 4, 5, 6, 7, 8, 9}
d_vec = thrust::sequence(d_vec.begin(), d_vec.begin() + 10);

//declare a custom operator
struct add_5{
    __host__ __device__ int operator()(int a){
        return a + 5;
    }
};

add_5 func;

//apply custom transformation
thrust::transform(d_vec.begin(), d_vec.end(), d_vec_out.begin(), func);

//d_vec is now {5, 6, 7, 8, 9, 10, 11, 12, 13, 14}
```



Thrust Fusion

❑ For best performance it is necessary to fuse operations



```
struct absolute{
    __host__ __device__ int operator()(int a){
        return a < 0 ? -a : a ;
    }
};
absolute func;

//custom transformation to calculate absolute value
thrust::transform(d_vec.begin(), d_vec.end(), d_vec.begin(), func);
//apply reduction, maximum binary associate operator
int result = thrust::reduce(d_vec.begin(), d_vec.end(), 0, thrust::maximum<int>());
```

```
struct absolute{
    __host__ __device__ int operator()(int a){
        return a < 0 ? -a : a ;
    }
};
absolute func;

//apply transform reduction maximum binary associate operator
int result = thrust::transform_reduce(d_vec.begin(), d_vec.end(), func, 0, thrust::maximum<int>());
```


Summary

❑ Libraries and Thrust

- ❑ Describe the purpose of CUDA libraries
- ❑ Demonstrate Thrust containers for data storage
- ❑ Explain the relationship between raw pointers and Thrust iterators
- ❑ Give example of Thrust algorithms

❑ Next: Applications of Sort

