

Parallel Computing with GPUs

CUDA Streams

Part 2 – CUDA Streams



Dr Paul Richmond

<http://paulrichmond.shef.ac.uk/teaching/COM4521/>



This Lecture (learning objectives)

□ CUDA Streams

- Demonstrate how to make asynchronous memory calls
- Demonstrate how have *copy* and *compute* concurrency
- Give examples of the issue ordering implications of stream scheduling and concurrency



Opportunities for Device Concurrency

- ❑ Most CUDA Devices have an asynchronous Kernel execution and Copy Engine
 - ❑ Allows data to be moved at the same time as execution
 - ❑ Most device have dual copy engines
 - ❑ PCIe upstream (D2H)
 - ❑ PCIe downstream (H2D)
 - ❑ Ideally we should hide data movement with execution
 - ❑ Check your device capability: `deviceQuery` example “Concurrent copy and kernel execution:
- ❑ All modern GPU devices are able to execute kernels simultaneously
 - ❑ Allows task parallelism on GPU
 - ❑ Each kernel represents a different task
 - ❑ Very useful for smaller problem sizes



Streams

- ❑ CUDA Streams allow operations to be queued for the GPU device
 - ❑ All calls are asynchronous by default
 - ❑ The host retains control
 - ❑ Device takes work from the streams when it is able to do so
- ❑ Operations in a stream are ordered and can not overlap (FIFO)
- ❑ Operations in different streams can overlap

```
// create a handle for the stream
cudaStream_t stream;
//create the stream
cudaStreamCreate(&stream);

//do some work in the stream ...

//destroy the stream (blocks host until stream is complete)
cudaStreamDestroy(stream);
```



Work Assignment for Streams



```
//execute kernel on device in specified stream  
fooKernel<<<blocks, threads, 0, stream>>>();
```

❑ Kernel Execution is assigned to streams as 4th parameter of kernel launch

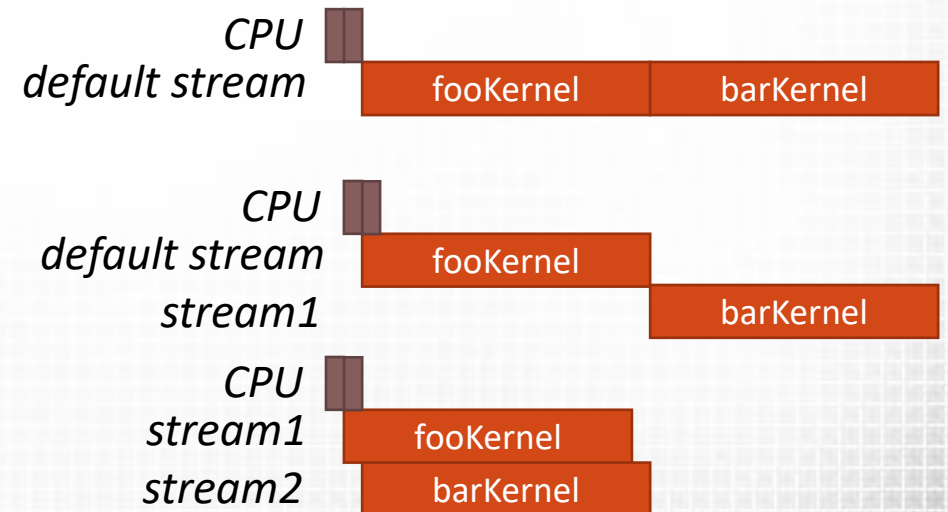
❑ Care must be taken with the default stream

❑ Only stream which is synchronous with others!

```
fooKernel<<<blocks, threads, 0>>>();  
barKernel<<<blocks, threads, 0>>>();
```

```
fooKernel<<<blocks, threads, 0>>>();  
barKernel<<<blocks, threads, 0, stream>>>();
```

```
fooKernel<<<blocks, threads, 0, stream1>>>();  
barKernel<<<blocks, threads, 0, stream2>>>();
```



Asynchronous Memory

- ❑ CUDA is able to asynchronously copy data
 - ❑ Only if it is Pinned (Page-locked) memory
- ❑ Paged Memory
 - ❑ Allocated using `malloc (...)` on host and released using `free (...)`
- ❑ Pinned Memory
 - ❑ Can not be swapped (paged) out by the OS
 - ❑ Has higher overhead for allocation
 - ❑ Can reach higher bandwidths for large transfers
 - ❑ Allocated using `cudaMallocHost (...)` and released using `cudaFreeHost (...)`
 - ❑ Can also pin non pinned memory using `cudaHostRegister (...)` / `cudaHostUnregister (...)`
 - ❑ Very slow



Concurrent Copies in Streams

- ❑ Memory copies can be replaced with `cudaMemcpyAsync()`
 - ❑ Requires an extra argument (a stream)
 - ❑ Places transfer into the stream and returns control to host
 - ❑ Conditions of use
 - ❑ Must be pinned memory
 - ❑ Must be in the non-default stream

```
int *h_A, *d_A;
cudaStream_t stream1;

cudaStreamCreate(&stream1);
cudaMallocHost(&h_A, SIZE);
cudaMalloc(&d_A, SIZE);
initialiseA(h_A);

cudaMemcpyAsync(d_A, h_A, SIZE, cudaMemcpyHostToDevice, stream1);

//work in other streams ...

cudaStreamDestroy(stream1);
```

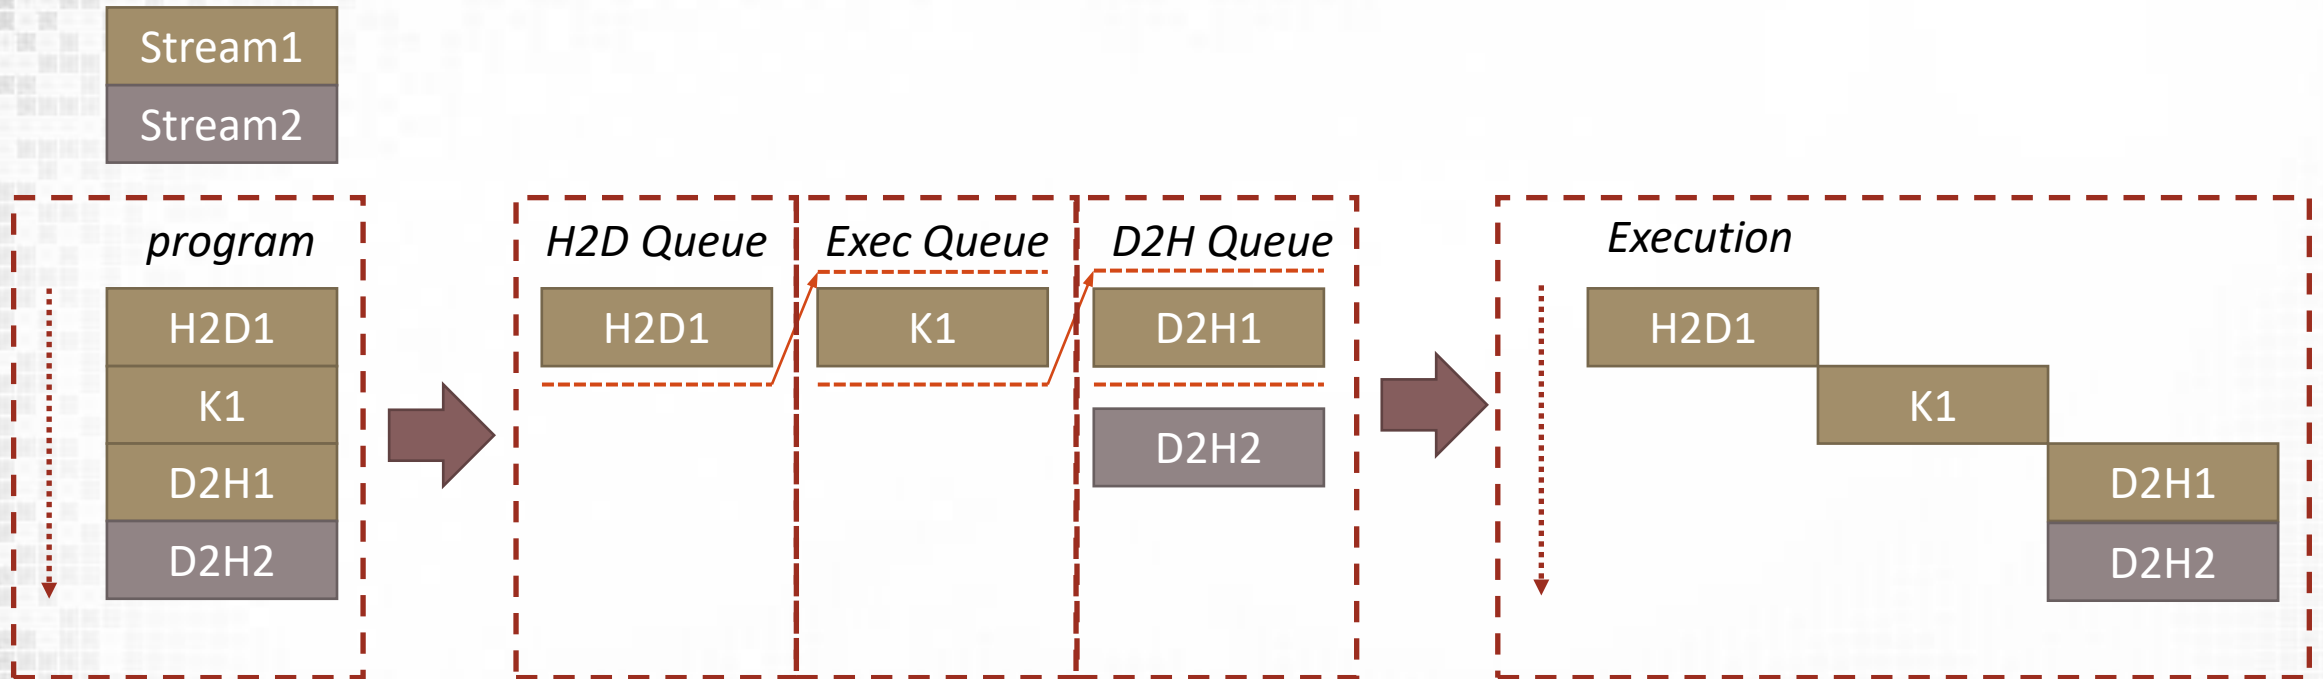


Stream Scheduling

- ❑ CUDA operations dispatched to hardware in sequence that they were issued
 - ❑ Hence issue order is important (FIFO)
- ❑ Kernel and Copy Engine (x2) have different queues
- ❑ Operations are de-queued if
 1. Preceding call in the same stream have completed
 2. Preceding calls in the same queue have been dispatched, and
 3. Resources are available
 - ❑ i.e. kernels can be concurrently executed if in different streams
- ❑ Blocking operations (e.g. `cudaMemcpy` will block all streams)



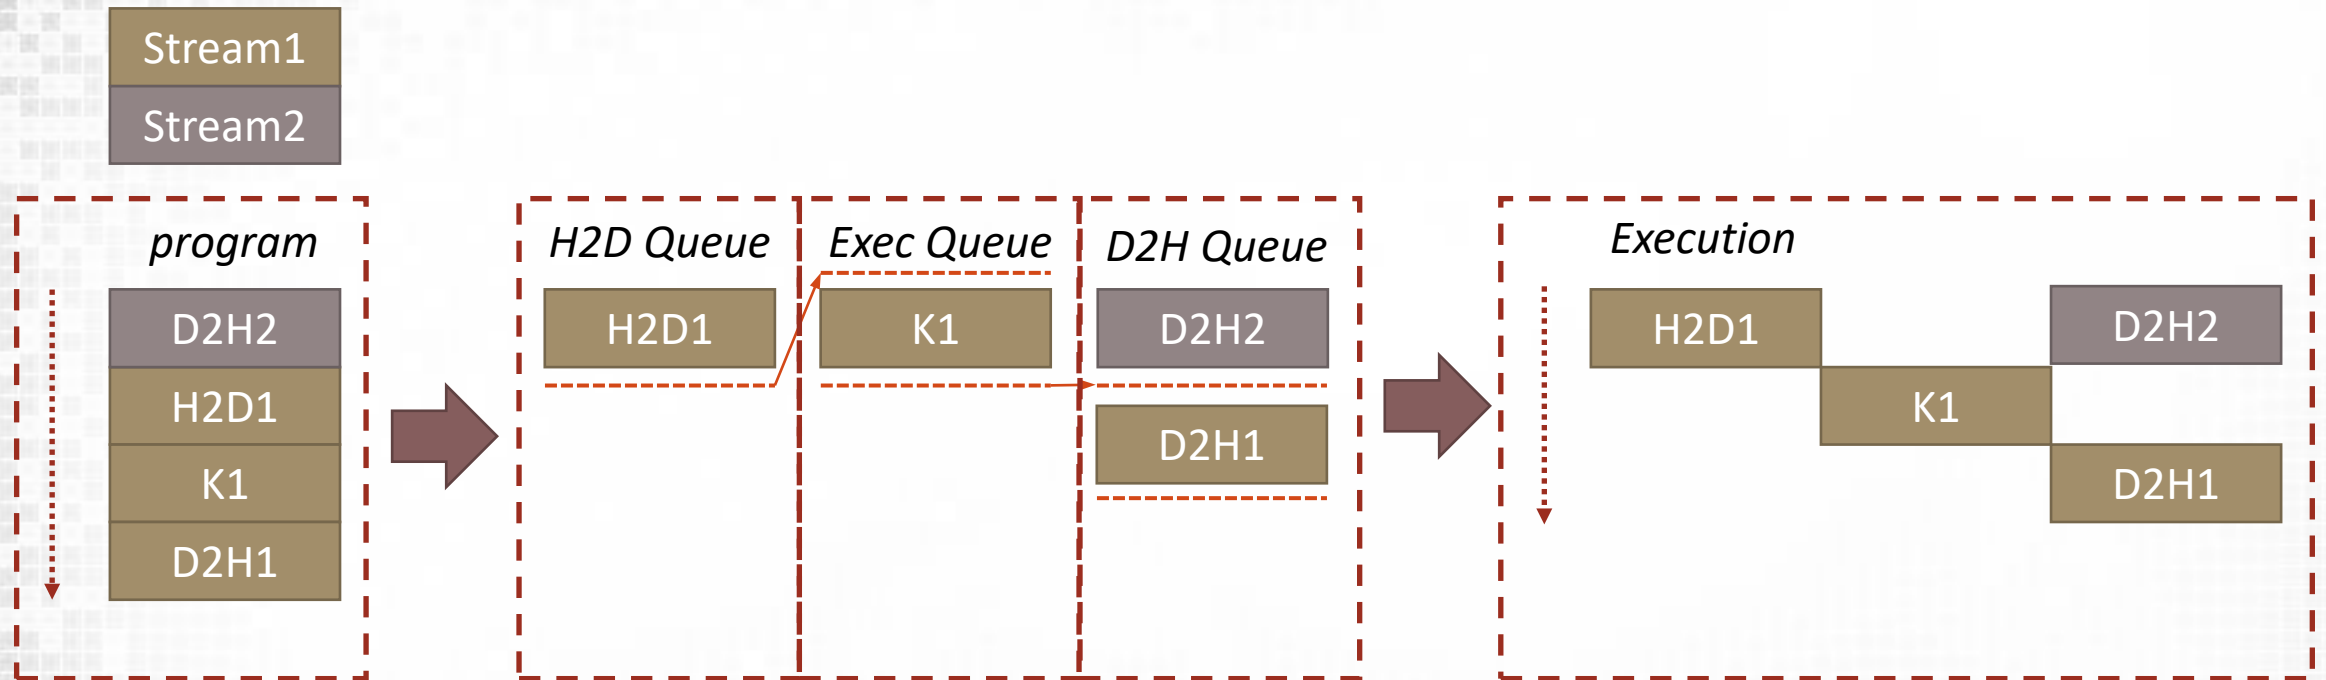
Issue Ordering



- ❑ No Concurrency of D2H2
- ❑ Blocked by D2H1
 - ❑ Issued first (FIFO)



Issue Ordering



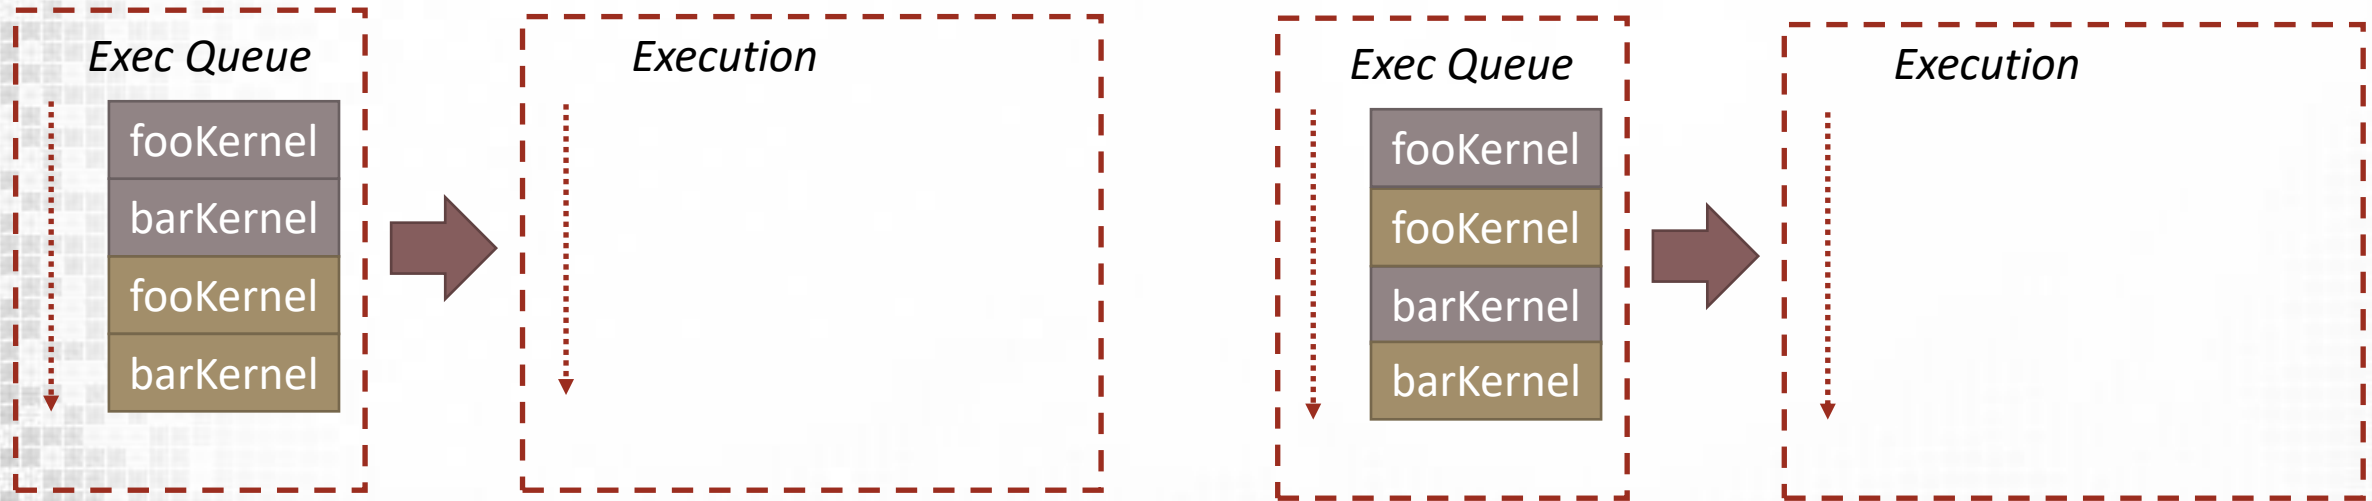
❑ Concurrency of D2H2 and H2D1





Issue Ordering (Kernel Execution)

Stream1
Stream2




❑ Which has best Asynchronous execution?



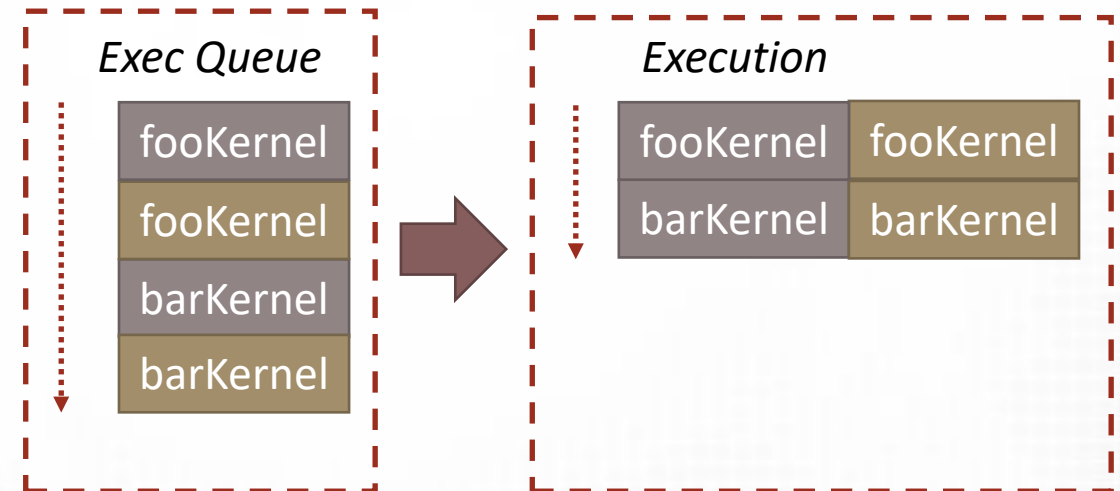
Issue Ordering (Kernel Execution)



Stream1
Stream2



- ☐ barKernel can't be removed from queue until fooKernel has completed
- ☐ Blocks **fooKernel**




- ☐ Both fooKernels can be concurrently executed
- ☐ Both barKernels concurrently executed

Levels of Concurrency




Fully Synchronous (Serial Execution)




2-way Concurrency

- ☐ H2D and D2H not concurrent



3-way Concurrency

- ☐ Both Copy Engines active
- ☐ Execution Engine active
 - ☐ May or may not be fully utilised



5-way Concurrency

- ☐ Both Copy Engines active
- ☐ Execution Engine active
 - ☐ Higher independent workload
 - ☐ Better chance of 100% utilisation
- ☐ What about Host?

Summary

❑ CUDA Streams

- ❑ Demonstrate how to make asynchronous memory calls
- ❑ Demonstrate how have *copy* and *compute* concurrency
- ❑ Give examples of the issue ordering implications of stream scheduling and concurrency

❑ Next Lecture: Synchronisation

