

Parallel Computing with GPUs

CUDA Streams

Part 3 – Synchronisation



The
University
Of
Sheffield.

Dr Paul Richmond

<http://paulrichmond.shef.ac.uk/teaching/COM4521/>



This Lecture (learning objectives)

☐ Synchronisation

- ☐ Identify and explain different levels of synchronisation
- ☐ Introduce and give examples of events
- ☐ Demonstrate the use of callbacks



Explicit Device Synchronisation

- ❑ What if we want to ensure an asynchronous kernel call has completed?
 - ❑ For timing kernel execution
 - ❑ Accessing data copied asynchronously without causing race conditions
- ❑ `cudaDeviceSynchronize()`
 - ❑ Will ensure that all asynchronous device operations are completed
 - ❑ Synchronise everything!
- ❑ `cudaStreamSynchronize(stream)`
 - ❑ Blocks host until all calls in stream are complete
- ❑ *CUDA Event synchronisation...*



Events

- ❑ Mechanism in which to signal when operations have occurred in a stream
 - ❑ Places an event into a stream (default stream unless specified)
- ❑ We have seen events already!
 - ❑ When timing our code...

```
cudaEvent_t start, stop;  
cudaEventCreate(&start);  
cudaEventCreate(&stop);  
  
cudaEventRecord(start);  
my_kernel <<<(N /TPB), TPB >>>();  
cudaEventRecord(stop);  
  
cudaEventSynchronize(stop);  
float milliseconds = 0;  
cudaEventElapsedTime(&milliseconds, start, stop);  
  
cudaEventDestroy(start);  
cudaEventDestroy(stop);
```



Events and Streams

- ❑ `cudaEventRecord(event, stream)`
 - ❑ Places an event in the non default stream
- ❑ `cudaEventSynchronize(event)`
 - ❑ Blocks until the stream completes all outstanding calls
 - ❑ Should be called after the event is inserted into the stream
- ❑ `cudaStreamWaitEvent(stream, event)`
 - ❑ Blocks the stream until the event occurs
 - ❑ Only blocks launches after event
 - ❑ Does not block the host
- ❑ `cudaEventQuery(event, stream)`
 - ❑ Has the event occurred in the stream

```
cudaMemcpyAsync(d_in, in, size, H2D, stream1);  
cudaEventRecord(event, stream1); // record event  
  
cudaStreamWaitEvent(stream2, event); // wait for event in stream1  
kernel << <BLOCKS, TPB, 0, stream2 >> > (d_in, d_out);
```



Callbacks

- ❑ Callbacks are functions on the host which should be called when an event is reached
- ❑ `cudaStreamAddCallback(stream, callback, user_data, 0)`
 - ❑ Good for scheduling host code once event has completed
 - ❑ Allows GPU to initiate operations that only the CPU can perform
 - ❑ Disk or network IO
 - ❑ System calls, etc.

```
void CUDART_CB MyCallback(void *data) {  
    //some host code  
}
```

```
MyKernel << <BLOCKS, TPB, 0, stream >> >(d_i);  
cudaStreamAddCallback(stream, MyCallback, (void*)d_i, 0);
```



WDDM Command Queues

- ❑ GPUs driving a display in windows use the Windows Display Driver Model Command Queues.
 - ❑ All CUDA calls (sync/async) are buffered within a WDDM *Command Buffer*
 - ❑ The *Command Buffer* will usually be flushed by
 - ❑ Forcing it by calling `cudaEventQuery(0)`
 - ❑ Issuing a synchronous call. E.g. a stream/device sync or synchronous memcpy
 - ❑ Waiting until it gets full (unpredictable)
 - ❑ Magic???
- ❑ Implications
 - ❑ Only things in the same command buffer can be concurrent
 - ❑ Stuff *might* not get queued into copy/compute engines as you would expect
 - ❑ `cudaEventElapsedTime` *may* not be accurate for asynchronous host timing



Summary

❑ Synchronisation

- ❑ Identify and explain different levels of synchronisation
- ❑ Introduce and give examples of events
- ❑ Demonstrate the use of callbacks

❑ Next Lecture: Multi GPU Programming

