

# Parallel Computing with GPUs

## CUDA Streams Part 4 – Multi GPU Programming



Dr Paul Richmond

<http://paulrichmond.shef.ac.uk/teaching/COM4521/>



# This Lecture (learning objectives)

## ❑ Multi GPU Programming

- ❑ Demonstrate how to change the GPU device
- ❑ Explain how devices can be concurrently operated and synchronised using streams
- ❑ Demonstrate mechanisms for device to device asynchronous memory copying



# Multi GPU Programming

- ❑ By default CUDA uses the first device in the system
  - ❑ Not necessarily the fastest device!
- ❑ Device can be changed using `cudaSetDevice(int)`
  - ❑ Device capabilities can be queried using device properties API

```
int deviceCount = 0;
cudaGetDeviceCount(&deviceCount);

for (int dev = 0; dev < deviceCount; ++dev)
{
    cudaSetDevice(dev);
    cudaDeviceProp deviceProp;
    cudaGetDeviceProperties(&deviceProp, dev);
    ...
}
```



# Multi GPU Devices and Streams

- ❑ Streams and events belong to a single device
  - ❑ The device which is active when created
  - ❑ Synchronising and Querying of streams across devices is allowed

```
cudaStream_t streamA, streamB;  
cudaEvent_t eventA, eventB;  
  
cudaSetDevice(0);  
cudaStreamCreate(&streamA); // streamA and eventA belong to device-0  
cudaEventCreate(&eventA);  
  
cudaSetDevice(1);  
cudaStreamCreate(&streamB); // streamB and eventB belong to device-1  
cudaEventCreate(&eventB);  
kernel << <..., streamB >> >(...);  
cudaEventRecord(eventB, streamB);  
  
cudaSetDevice(0);  
cudaEventSynchronize(eventB);  
kernel << <..., streamA >> >(...);
```



Error: eventA belongs to device 0

Event can be synchronised across devices



# Peer to Peer Memory Copies

- ❑ For devices to interact memory must be copied between them
- ❑ Memory can be copied using
  - ❑ `cudaMemcpyPeerAsync( void* dst_addr, int dst_dev, void* src_addr, int src_dev, size_t num_bytes, cudaStream_t stream )`
    - ❑ Uses shortest PCI path or GPUDirect if available
    - ❑ Not staged through CPU
- ❑ You can check that a peer (device) can access another using
  - ❑ `cudaDeviceCanAccessPeer( &accessible, dev_X, dev_Y )`
- ❑ Also possible to use CUDA aware MPI
  - ❑ Allows direct transfers over the network
  - ❑ With NVLink this will allow GPU to GPU peer access via infiniband
  - ❑ *Not covered in this course...*



# Further Reading & Acknowledgements

- ❑ Most slide examples are based on the excellent GTC and SC material
  - ❑ <http://www.sie.es/wp-content/uploads/2015/12/cuda-streams-best-practices-common-pitfalls.pdf>
  - ❑ <http://on-demand.gputechconf.com/gtc-express/2011/presentations/StreamsAndConcurrencyWebinar.pdf>
  - ❑ <http://www.nvidia.com/docs/IO/116711/sc11-multi-gpu.pdf>
- ❑ More reading
  - ❑ <https://devblogs.nvidia.com/parallelforall/gpu-pro-tip-cuda-7-streams-simplify-concurrency/>
  - ❑ <https://devblogs.nvidia.com/parallelforall/how-overlap-data-transfers-cuda-cc/>



# Summary

## ❑ Multi GPU Programming

- ❑ Demonstrate how to change the GPU device
- ❑ Explain how devices can be concurrently operated and synchronised using streams
- ❑ Demonstrate mechanisms for device to device asynchronous memory copying

